

FTI - Fault Tolerance Interface

Contact:

Dr. Leonardo Bautista-Gomez (Leo) - leonardo.bautista@bsc.es

Barcelona Supercomputing Center

Carrer de Jordi Girona, 29-31, 08034 Barcelona, SPAIN

Phone : +34 934 13 77 16



Authors: Kai Keller, Tomasz Paluszkiwicz, Karol Sierocinski

Release: 0.9.9

In high performance computing (HPC), systems are built from highly reliable components. However, the overall failure rate of supercomputers increases with component count. Nowadays, petascale machines have a mean time between failures (MTBF) measured in hours or days and fault tolerance (FT) is a well-known issue. Long running large applications rely on FT techniques to successfully finish their long executions. Checkpoint/Restart (CR) is a popular technique in which the applications save their state in stable storage, frequently a parallel file system (PFS); upon a failure, the application restarts from the last saved checkpoint. CR is a relatively inexpensive technique in comparison with the process-replication scheme that imposes over 100% of overhead.

However, when a large application is checkpointed, tens of thousands of processes will each write several GBs of data and the total checkpoint size will be in the order of several tens of TBs. Since the I/O bandwidth of supercomputers does not increase at the same speed as computational capabilities, large checkpoints can lead to an I/O bottleneck, which causes up to 25% of overhead in current petascale systems. Post-petascale systems will have a significantly larger number of components and an important amount of memory. This will have an impact on the system's reliability. With a shorter MTBF, those systems may require a higher checkpoint frequency and at the same time they will have significantly larger amounts of data to save. Although the overall failure rate of future post-petascale systems is a common factor to study when designing FT-techniques, another important point to take into account is the pattern of the failures. Indeed, when moving from 90nm to 16nm technology, the soft error rate (SER) is likely to increase significantly, as shown in a recent study from Intel. A recent study by Dong et al. explains how this provides an opportunity for local/global hybrid checkpoint using new technologies such as phase change memories (PCM). Moreover, some hard failures can be tolerated using solid-state-drives (SSD) and cross-node redundancy schemes, such as checkpoint replication or XOR encoding which allows to leverage multi-level checkpointing, as proposed by Moody et al.. Furthermore, Cheng et al. demonstrated that more complex erasure codes such as Reed-Solomon (RS) encoding can be used to further increase the percentage of hard failures tolerated without stressing the PFS.

FTI is a multi-level checkpointing interface. It provides an api which is easy to apply and offers a flexible configuration to enable the user to select the checkpointing strategy which fits best to the problem.

L1

L1 denotes the first safety level in the multilevel checkpointing strategy of FTI. The checkpoint of each process is written on the local SSD of the respective node. This is fast but possesses the drawback, that in case of a data loss and corrupted checkpoint data even in only one node, the execution cannot successfully restarted.

L2

L2 denotes the second safety level of checkpointing. On initialisation, FTI creates a virtual ring for each group of nodes with user defined size (see [group_size](#)). The first step of L2 is just a L1 checkpoint. In the second step, the checkpoints are duplicated and the copies stored on the neighbouring node in the group.

That means, in case of a failure and data loss in the nodes, the execution still can be successfully restarted, as long as the data loss does not happen on two neighbouring nodes at the same time.

L3

L3 denotes the third safety level of checkpointing. In this level, the checkpoint data trunks from each node getting encoded via the Reed-Solomon (RS) erasure code. The implementation in FTI can tolerate the breakdown and data loss in half of the nodes.

In contrast to the safety level L2, in level L3 it is irrelevant which of nodes encounters the failure. The missing data can get reconstructed from the remaining RS-encoded data files.

L4

L4 denotes the fourth safety level of checkpointing. All the checkpoint files are flushed to the parallel file system (PFS).

FTI uses Cmake to configure the installation. The recommended way to perform the installation is to create a build directory within the base directory of FTI and perform the cmake command in there. In the following you will find configuration examples. The commands are performed in the build directory within the FTI base directory.

Default The default configuration builds the FTI library with Fortran and MPI-IO support for GNU compilers:

```
cmake -DCMAKE_INSTALL_PREFIX:PATH=/install/here/fti ..  
make all install
```

Notice: THE TWO DOTS AT THE END INVOKE CMAKE IN THE TOP LEVEL DIRECTORY.

Intel compilers Fortran and MPI-IO support for Intel compilers:

```
cmake -C ../intel.cmake -DCMAKE_INSTALL_PREFIX:PATH=/install/here/fti ..  
make all install
```

Disable Fortran Only build FTI C library:

```
cmake -DCMAKE_INSTALL_PREFIX:PATH=/install/here/fti -DENABLE_FORTRAN=OFF ..  
make all install
```

Lustre For Lustre user who want to use MPI-IO, it is strongly recommended to configure with Lustre support:

```
cmake -DCMAKE_INSTALL_PREFIX:PATH=/install/here/fti -DENABLE_LUSTRE=ON ..  
make all install
```

Cray For Cray systems, make sure that the modules craype/* and PrgEnv* are loaded (if available). The configuration should be done as:

```
export CRAY_CPU_TARGET=x86-64  
export CRAYPE_LINK_TYPE=dynamic  
cmake -DCMAKE_INSTALL_PREFIX:PATH=/install/here/fti -DCMAKE_SYSTEM_NAME=CrayLinuxEnvironment ..  
make all install
```

Notice: MODIFY x86-64 IF YOU ARE USING A DIFFERENT ARCHITECTURE. ALSO, THE OPTION `CMAKE_SYSTEM_NAME=CrayLinuxEnvironment` IS AVAILABLE ONLY FOR CMAKE VERSIONS 3.5.2 AND ABOVE.

FTI Datatypes and Constants

FTI Datatypes

`FTI_CHAR` : FTI data type for chars
`FTI_SHRT` : FTI data type for short integers.
`FTI_INTG` : FTI data type for integers.
`FTI_LONG` : FTI data type for long integers.
`FTI_UCHR` : FTI data type for unsigned chars.
`FTI_USHT` : FTI data type for unsigned short integers.
`FTI_UINT` : FTI data type for unsigned integers.
`FTI_ULNG` : FTI data type for unsigned long integers.
`FTI_SFLT` : FTI data type for single floating point.
`FTI_DBLE` : FTI data type for double floating point.
`FTI_LDBE` : FTI data type for long double floating point.

FTI Constants

`FTI_BUFS` : 256
`FTI_DONE` : 1
`FTI_SCES` : 0
`FTI_NSCS` : -1
`FTI_NREC` : -2

FTI_Init

- Reads configuration file.
- Creates checkpoint directories.
- Detects topology of the system.
- Regenerates data upon recovery.

DEFINITION

```
int FTI_Init ( char * configFile , MPI_Comm globalComm )
```

INPUT

Variable	What for?
<code>char * configFile</code>	Path to the config file
<code>MPI_Comm globalComm</code>	MPI communicator used for the execution

OUTPUT

Value	Reason
<code>FTI_SCES</code>	Success
<code>FTI_NSCS</code>	No Success
<code>FTI_NREC</code>	FTI could not recover ckpt files

DESCRIPTION

`FTI_Init` initializes the FTI context. It must be called before any other FTI function and after `MPI_Init`.

EXAMPLE

```

int main ( int argc , char **argv ) {
    MPI_Init (&argc , &argv );
    char *path = "config.fti"; // config file path
    int res = FTI_Init ( path , MPI_COMM_WORLD );
    if (res == FTI_NREC) {
        printf("Recovery not possible, terminating...");
        FTI_Finalize();
        MPI_Finalize();
        return 1;
    }
    .
    .
    .
    return 0;
}

```

FTI_InitType

- Initializes a data type.

DEFINITION

```
int FTI_InitType ( FTIT_type *type , int size )
```

INPUT

Variable	What for?
<code>FTIT_type * type</code>	The data-type to be initialized
<code>int size</code>	The size of the data-type to be initialized

OUTPUT

Value	Reason
<code>FTI_SCES</code>	Success

DESCRIPTION

FTI_InitType initializes a FTI data-type. A data-type which is not defined by default by FTI (see [FTI Datatypes](#)), must be defined using this function in order to protect variables of that type with [FTI_Protect](#).

EXAMPLE

```

typedef struct A {
    int a;
    int b;
} A;
FTIT_type structAinfo ;
FTI_InitType (& structAinfo , 2 * sizeof ( int ));

```

FTI_Protect

- Stores metadata concerning the variable to protect.

DEFINITION

```
int FTI_Protect ( int id, void *ptr, long count, FTIT_type type )
```

INPUT

Variable	What for?
<code>int id</code>	Unique ID of the variable to protect
<code>void * ptr</code>	Pointer to memory address of variable
<code>long count</code>	Number of elements at memory address
<code>FTIT_type type</code>	FTI data type of variable to protect

OUTPUT

Value	Reason
<code>FTI_SCES</code>	Success
<code>FTI_NSCS</code>	No success

DESCRIPTION

`FTI_Protect` is used to add data fields to the list of protected variables. Data, protected by this function will be stored during a call to [FTI_Checkpoint](#) or [FTI_Snapshot](#) and restored during a call to [FTI_Recover](#).

If the dimension of a protected variable changes during the execution, a subsequent call to `FTI_Protect` will update the meta-data within FTI in order to store the correct size during a successive call to [FTI_Checkpoint](#) or [FTI_Snapshot](#).

EXAMPLE

```
int A;
float *B = malloc (sizeof(float) * 10) ;
FTI_Protect(1, &A, 1, FTI_INTG );
FTI_Protect(2, B, 10, FTI_SFLT );
// changing B size
B = realloc(B, sizeof(float) * 20) ;
// updating B size in protected list
FTI_Protect(2, B, 20, FTI_SFLT);
```

FTI_GetStoredSize

- Returns size of protected variable saved in metadata

DEFINITION

```
long FTI_GetStoredSize ( int id )
```

INPUT

Variable	What for?
<code>int id</code>	ID of the protected variable

OUTPUT

Value	Reason
<code>long</code>	Size of a variable
<code>0</code>	No success

DESCRIPTION

`FTI_GetStoredSize` returns the size of a protected variable with `id` from the FTI metadata. The result may differ from the size of the variable known to the application at that moment. If the function is called on a restart, it returns the size stored in the

metadata file. Called during the execution, it returns the value stored in the FTI runtime metadata, i.e. the size of the variable at the moment of the last checkpoint.

The function is needed to manually reallocate memory for protected variables with variable size on a recovery. Another possibility for the reallocation of memory is provided by [FTI_Realloc](#).

EXAMPLE

```
...
long* array = calloc(arraySize, sizeof(long));
FTI_Protect(1, array, arraySize, FTI_LONG);
if (FTI_Status() != 0) {
    long arraySizeInBytes = FTI_GetStoredSize(1);
    if (arraySizeInBytes == 0) {
        printf("No stored size in metadata!\n");
        return GETSTOREDSIZE_FAILED;
    }
    array = realloc(array, arraySizeInBytes);
    int res = FTI_Recover();
    if (res != 0) {
        printf("Recovery failed!\n");
        return RECOVERY_FAILED;
    }
    //update arraySize
    arraySize = arraySizeInBytes / sizeof(long);
}
for (i = 0; i < max; i++) {
    if (i % CKTP_STEP) {
        //update FTI array size information
        FTI_Protect(1, array, arraySize, FTI_LONG);
        int res = FTI_Checkpoint((i % CKTP_STEP) + 1, 1);
        if (res != FTI_DONE) {
            printf("Checkpoint failed!\n");
            return CHECKPOINT_FAILED;
        }
    }
}
...
//add element to array
arraySize += 1;
array = realloc(array, arraySize * sizeof(long));
}
...
```

FTI_Realloc

Reallocates dataset to last checkpoint size.

DEFINITION

```
void* FTI_Realloc ( int id, void* ptr )
```

INPUT

Variable	What for?
int id	ID of the protected variable
void * ptr	Pointer to memory address of variable

OUTPUT

Value	Reason
void*	Pointer to reallocated data
NULL	On failure

DESCRIPTION

`FTI_Realloc` is called for protected variables with dynamic size on recovery. It reallocates sufficient memory to store the checkpoint data to the pointed memory address. It must be called before `FTI_Recover` to prevent segmentation faults. If the reallocation must/is wanted to be done within the application, FTI provides the function `FTI_GetStoredSize` to request the variable size of the checkpoint to recover.

EXAMPLE

```
...
FTI_Protect(1, &arraySize, 1, FTI_INTG);
long* array = calloc(arraySize, sizeof(long));
FTI_Protect(2, array, arraySize, FTI_LONG);
if (FTI_Status() != 0) {
    array = FTI_Realloc(2, array);
    if (array == NULL) {
        printf("Reallocation failed!\n");
        return REALLOC_FAILED;
    }

    int res = FTI_Recover();
    if (res != 0) {
        printf("Recovery failed!\n");
        return RECOVERY_FAILED;
    }
}
for (i = 0; i < max; i++) {
    if (i % CKTP_STEP) {
        //update FTI array size information
        FTI_Protect(2, array, arraySize, FTI_LONG);
        int res = FTI_Checkpoint((i % CKTP_STEP) + 1, 1);
        if (res != FTI_DONE) {
            printf("Checkpoint failed!\n");
            return CHECKPOINT_FAILED;
        }
    }
}
...
//add element to array
arraySize += 1;
array = realloc(array, arraySize * sizeof(long));
}
...
```

FTI_Checkpoint

- Stores protected variables in the checkpoint of a desired safety level.

DEFINITION

```
int FTI_Checkpoint( int id, int level )
```

INPUT

Variable	What for?
<code>int id</code>	Unique checkpoint ID
<code>int level</code>	Checkpoint level (1=L1, 2=L2, 3=L3, 4=L4)

OUTPUT

Value	Reason
<code>FTI_DONE</code>	Success
<code>FTI_NSCS</code>	Failure

DESCRIPTION

FTI_Checkpoint is used to store the current values of protected variables into a checkpoint of safety level `level` (see [Multilevel-Checkpointing](#) for descriptions of the particular levels).

NOTICE: The checkpoint id must be different from 0!

EXAMPLE

```
int i;
for (i = 0; i < 100; i++) {
    if (i % 10 == 0) {
        FTI_Checkpoint ( i/10 + 1, 1 );
    }
    .
    . // some computations
    .
}
```

FTI_Status

- Returns the current status of the recovery flag.

DEFINITION

```
int FTI_Status()
```

OUTPUT

Value	Reason
int 0	No checkpoints taken yet or recovered successfully
int 1	At least one checkpoint is taken. If execution fails, the next start will be a restart
int 2	The execution is a restart from checkpoint level L4 and <code>keep_last_checkpoint</code> was enabled during the last execution

DESCRIPTION

FTI_Status returns the current status of the recovery flag.

EXAMPLE

```
if ( FTI_Status () != 0 ) {
    .
    . // this section will be executed during restart
    .
}
```

FTI_Recover

- Recovers the data of the protected variables from the checkpoint file.

DEFINITION

```
int FTI_Recover()
```

OUTPUT

Value	Reason
FTI_SCES	Success
FTI_NSCS	Failure

DESCRIPTION

`FTI_Recover` loads the data from the checkpoint file to the protected variables. It only recovers variables which are protected by a preceding call to `FTI_Protect`. If a variable changes its size during execution, the proper amount of memory has to be allocated for that variable before the call to `FTI_Recover`. FTI provides the API functions `FTI_GetStoredSize` and

EXAMPLE

Basic example:

```
if ( FTI_Status() == 1 ) {
    FTI_Recover();
}
```

FTI_Snapshot

- Invokes the recovery of protected variables on a restart.
- Writes multilevel checkpoints regarding their requested frequencies during execution.

DEFINITION

```
int FTI_Snapshot()
```

OUTPUT

Value	Reason
FTI_SCES	Successfull call (without checkpointing) or if recovery successful
FTI_NSCS	Failure of <code>FTI_Checkpoint</code>
FTI_DONE	Success of <code>FTI_Checkpoint</code>
FTI_NREC	Failure on recovery

DESCRIPTION

On a restart, `FTI_Snapshot` loads the data from the checkpoint file to the protected variables. During execution it performs checkpoints according to the checkpoint frequencies for the various safety levels. The frequencies may be set in the configuration file (see e.g.: [ckpt_L1](#)).

`FTI_Snapshot` can only take care of variables which are protected by a preceding call to `FTI_Protect`.

EXAMPLE

```
int res = FTI_Snapshot();
if ( res == FTI_SCES ) {
    .
    . // executed after successful recover
    . // or when checkpoint is not required
}
else { // res == FTI_DONE
    .
    . // executed after successful checkpointing
    .
}
```

FTI_Finalize

- Frees the allocated memory.
- Communicates the end of the execution to dedicated threads.
- Cleans checkpoints and metadata.

DEFINITION

```
int FTI_Finalize()
```

OUTPUT

Value	Reason
FTI_SCES	For application process
exit(0)	For FTI process

DESCRIPTION

FTI_Finalize notifies the FTI processes that the execution is over, frees FTI internal data structures and it performs a clean up of the checkpoint folders at a normal execution. If the setting keep_last_ckpt is set, it flushes local checkpoint files (if present) to the PFS. If the setting head is set to 1, it will also terminate the FTI processes. It should be called before MPI_Finalize() .

EXAMPLE

```
int main ( int argc , char ** argv ) {  
    .  
    .  
    .  
    FTI_Finalize () ;  
    MPI_Finalize () ;  
    return 0;  
}
```

FTI Datatypes and Constants

FTI Datatypes

FTI datatypes are used in the C-API function `FTI_Protect`. With the `count` parameter and the datatype, FTI is able to determine the size of the allocated memory region at `ptr`.

The FTI Fortran interface defines a template of `FTI_Protect` for all intrinsic data types. Hence the datatype definitions are not necessary here and are not available for the Fortran interface.

FTI Constants

`FTI_BUFS` : 256

`FTI_DONE` : 1

`FTI_SCES` : 0

`FTI_NSCS` : -1

`FTI_NREC` : -2

FTI_Init

- Reads configuration file.
- Creates checkpoint directories.
- Detects topology of the system.
- Regenerates data upon recovery.

DEFINITION

```
subroutine FTI_Init ( config_file, global_comm, err )
```

ARGUMENTS

Variable		What for?
<code>character config_file</code>	IN	Path to the config file
<code>integer global_comm</code>	IN/OUT	MPI communicator used for the execution
<code>integer err</code>	OUT	Token for FTI error code.

ERROR HANDLING

ierr	Reason
<code>FTI_SCES</code>	Success
<code>FTI_NSCS</code>	No Success
<code>FTI_NREC</code>	FTI could not recover ckpt files

DESCRIPTION

`FTI_Init` initializes the FTI context. It must be called before any other FTI function and after `MPI_Init`. The MPI communicator passed, must be declared as `integer, target`.

EXAMPLE

```
integer, target :: rank, nbProcs, err, FTI_comm_world

call MPI_Init(err)
FTI_comm_world = MPI_COMM_WORLD
call FTI_Init('config.fti', FTI_comm_world, err) ! modifies FTI_comm_world
call MPI_Comm_size(FTI_comm_world, nbProcs, err)
call MPI_Comm_rank(FTI_comm_world, rank, err)
```

FTI_InitType

- Initializes a data type.

DEFINITION

```
subroutine FTI_InitType ( type_F, size_F, err )
```

ARGUMENTS

Variable		What for?
<code>type(FTI_type) type_F</code>	IN	The data type to be initialized
<code>integer size_F</code>	IN	The size of the data type to be initialized
<code>integer err</code>	OUT	Token for FTI error code.

ERROR HANDLING

err	Reason
<code>FTI_SCES</code>	Success
<code>FTI_NSCS</code>	No Success

DESCRIPTION

FTI_InitType initializes a FTI data-type. A data-type which is not Fortran intrinsic, must be defined using this function in order to protect variables of that type with [FTI_Protect](#).

EXAMPLE

```
!...

type polar
  real :: radius
  real :: phi
end type

type(FTI_Type)      :: FTI_Polar

type(polar), target :: choord
type(polar), pointer :: choord_ptr
type(c_ptr)         :: choord_c_ptr

choord_ptr => choord
choord_c_ptr = c_loc(choord)

! ...

call FTI_InitType(FTI_Polar, int(sizeof(choord),4), ierr)

! ...
```

FTI_Protect

- Stores metadata concerning the variable to protect.

In the Fortran interface, `FTI_Protect` comes with two different function headers. One may be used for intrinsic Fortran types and the other must be used for derived data-types.

DEFINITION

```
subroutine FTI_Protect ( id, data, err ) !> For intrinsic data-types
subroutine FTI_Protect ( id, data_ptr, count_F, type_F, err ) !> For derived data-types
```

ARGUMENTS (intrinsic types)

Variable		What for?
<code>integer id</code>	IN	Unique ID of the variable to protect
<code>Fortran type, pointer data</code>	IN	Pointer to memory address of variable
<code>integer err</code>	OUT	Token for FTI error code.

ARGUMENTS (derived types)

Variable		What for?
<code>integer id</code>	IN	Unique ID of the variable to protect
<code>type(c_ptr) data_ptr</code>	IN	Pointer to memory address of variable
<code>integer count_F</code>	IN	Number of elements.
<code>tape(FTI_Type) type_F</code>	IN	FTI_Type of Derived data-type.
<code>integer err</code>	OUT	Token for FTI error code.

ERROR HANDLING

err	Reason
<code>FTI_SCES</code>	Success
<code>FTI_NSCS</code>	Number of protected variables is > <code>FTI_BUFS</code>

DESCRIPTION

`FTI_Protect` is used to add data fields to the list of protected variables. Data, protected by this function will be stored during a call to [FTI_Checkpoint](#) or [FTI_Snapshot](#) and restored during a call to [FTI_Recover](#).

If the dimension of a protected variable changes during the execution, a subsequent call to `FTI_Protect` will update the meta-data within FTI in order to store the correct size during a successive call to [FTI_Checkpoint](#) or [FTI_Snapshot](#).

EXAMPLE

For Fortran intrinsic data-types:

```

! ...

integer, target :: nbProcs, iter, row, col, err, FTI_comm_world
integer, pointer :: ptriter
real(8), pointer :: g(:, :)

call MPI_Init(err)
FTI_comm_world = MPI_COMM_WORLD
call FTI_Init('config.fti', FTI_comm_world, err) ! modifies FTI_comm_world
call MPI_Comm_size(FTI_comm_world, nbProcs, err)

row = sqrt((MEM_MB * 1024.0 * 512.0 * nbProcs)/8)

col = (row / nbProcs)+3

allocate( g(row, col) )
allocate( h(row, col) )

! INIT DATA ...

ptriter => iter
call FTI_Protect(0, ptriter, err)
call FTI_Protect(2, g, err)

! ...

```

For derived data-types

```

! ...

use iso_c_binding

type polar
  real :: radius
  real :: phi
end type

type(FTI_Type)      :: FTI_Polar
integer, parameter  :: N=128*1024*25 !> 25 MB / Process
integer, parameter  :: N1 = 128
integer, parameter  :: N2 = 1024
integer, parameter  :: N3 = 25
integer, target     :: FTI_COMM_WORLD
integer             :: ierr, status

type(polar), dimension(:, :, :), pointer :: arr
type(c_ptr)        :: arr_c_ptr

allocate(arr(N1,N2,N3))

shape = (/ N1, N2, N3 /)
arr_c_ptr = c_loc( arr( &
  lbound(arr,1), &
  lbound(arr,2), &
  lbound(arr,3)))

!> INITIALIZE MPI AND FTI
call MPI_Init(ierr)
FTI_COMM_WORLD = MPI_COMM_WORLD
call FTI_Init('config.fti', FTI_COMM_WORLD, ierr)

call FTI_InitType(FTI_Polar, int(2*sizeof(1.0),4), ierr)

!> PROTECT DATA AND ITS SHAPE
call FTI_Protect(0, arr_c_ptr, size(arr), FTI_Polar, ierr)

! ...

```

FTI_Checkpoint

- Writes values of protected runtime variables to a checkpoint file of requested level.

DEFINITION

```
subroutine FTI_Checkpoint ( id_F, level, err )
```

ARGUMENTS

Variable		What for?
integer id_F	IN	Unique checkpoint ID
integer level	IN	Checkpoint level (1=L1, 2=L2, 3=L3, 4=L4)
integer err	OUT	Token for FTI error code.

ERROR HANDLING

err	Reason
FTI_DONE	Success
FTI_NSCS	Failure

DESCRIPTION

FTI_Checkpoint is used to store the current values of protected variables into a checkpoint of safety levellevel (see [Multilevel-Checkpointing](#) for descriptions of the particular levels).

EXAMPLE

The handling is identical to the C case, except that in Fortran it is *asubroutine* and not a function, hence:

```
! ...

!> LEVEL 2 CHECKPOINT, ID = 1
call FTI_Checkpoint(1, 2, err)

! ...
```

FTI_GetStoredSize

- Delivers the variable size in Bytes of a protected variable. The returned size is consistent to the FTI state, i.e. it might differ to the current variable size in the execution.

DEFINITION

```
subroutine FTI_GetStoredSize ( id_F, size_F )
```

ARGUMENTS

Variable		What for?
integer id_F	IN	Unique variable ID
integer size_F	OUT	Size of protected variable

ERROR HANDLING

size	Reason
> 0	Success
0	No size saved

DESCRIPTION

`FTI_GetStoredSize` returns the size of a protected variable with `id` from the FTI metadata. The result may differ from the size of the variable known to the application at that moment. If the function is called on a restart, it returns the size stored in the metadata file. Called during the execution, it returns the value stored in the FTI runtime metadata, i.e. the size of the variable at the moment of the last checkpoint.

The function is needed to manually reallocate memory for protected variables with variable size on a recovery. Another possibility for the reallocation of memory is provided by [FTI_Realloc](#).

EXAMPLE

```

! ...

integer, parameter :: N1=128*1024*25 !> 25 MB / Process
integer, parameter :: N2=128*1024*50 !> 50 MB / Process
integer :: varSizeMeta
integer, target :: FTI_COMM_WORLD
integer :: ierr, status

real(dp), dimension(:), pointer :: arr
real(dp), dimension(:), pointer :: tmp

allocate(arr(N1))

!> INITIALIZE MPI AND FTI
call MPI_Init(ierr)
FTI_COMM_WORLD = MPI_COMM_WORLD
call FTI_Init('config.fti', FTI_COMM_WORLD, ierr)

!> PROTECT DATA AND ITS SHAPE
call FTI_Protect(0, arr, ierr)

call FTI_Status(status)

!> EXECUTE ON RESTART
if ( status .eq. 1 ) then
  !> REALLOCATE TO SIZE AT CHECKPOINT
  call FTI_GetStoredSize(0, varSizeMeta)
  if(varSizeMeta .ne. sizeof(arr)) then
    deallocate(arr)
    allocate(arr(varSizeMeta))
    call FTI_Protect(0, arr, ierr) ! necessary to pass new address
  end if
  call MPI_Barrier(FTI_COMM_WORLD, ierr)
  call FTI_recover(ierr)

! ...

end if

! ...

!> FIRST CHECKPOINT
call FTI_Checkpoint(1, 1, ierr)

! ...

!> CHANGE ARRAY DIMENSION
allocate(tmp(N2))
tmp(1:N1) = arr
deallocate(arr)
arr => tmp

!> TELL FTI ABOUT THE NEW DIMENSION
call FTI_Protect(0, arr, ierr)

! ...

!> SECOND CHECKPOINT
call FTI_Checkpoint(2, 1, ierr)

! ...

```

FTI_Realloc

- Provides the reallocation of memory on FTI API side for protected variables upon a restart.

DEFINITION

```
subroutine FTL_Realloc ( id, data, err ) !> For intrinsic data-types
subroutine FTL_Realloc ( id, data_ptr, err ) !> For derived data-types
```

ARGUMENTS (intrinsic types)

Variable		What for?
integer id	IN	Unique ID of the variable to protect
Fortran type, pointer data	IN/OUT	Pointer to memory address of variable
integer err	OUT	Token for FTI error code.

ARGUMENTS (derived types)

Variable		What for?
integer id	IN	Unique ID of the variable to protect
type(c_ptr) data_ptr	IN/OUT	Pointer to memory address of variable
integer err	OUT	Token for FTI error code.

ERROR HANDLING

err	Reason
FTI_SCES	Success
FTI_NSCS	No success

DESCRIPTION

FTL_Realloc is called for protected variables with dynamic size on recovery. It reallocates sufficient memory to store the checkpoint data to the pointed memory address. It must be called before [FTI_Recover](#) to prevent segmentation faults. If the reallocation must/is wanted to be done within the application, FTI provides the function [FTI_GetStoredSize](#) to request the variable size of the checkpoint to recover.

EXAMPLE

For intrinsic data-types:

```

! ...

integer, parameter :: N1=128*1024*25 !> 25 MB / Process
integer, parameter :: N2=128*1024*50 !> 50 MB / Process
integer, parameter :: N11 = 128
integer, parameter :: N12 = 1024
integer, parameter :: N13 = 25
integer, parameter :: N21 = 128
integer, parameter :: N22 = 1024
integer, parameter :: N23 = 50
integer, target    :: FTI_COMM_WORLD
integer            :: ierr, status

real(dp), dimension(:, :, :), pointer :: arr
type(c_ptr)      :: arr_c_ptr
real(dp), dimension(:, :, :), pointer :: tmp
integer, dimension(:), pointer :: shape

allocate(arr(N11,N12,N13))
allocate(shape(3))

!> INITIALIZE MPI AND FTI
call MPI_Init(ierr)
FTI_COMM_WORLD = MPI_COMM_WORLD
call FTI_Init('config.fti', FTI_COMM_WORLD, ierr)

!> PROTECT DATA AND ITS SHAPE
call FTI_Protect(0, arr, ierr)
call FTI_Protect(1, shape, ierr)

call FTI_Status(status)

!> EXECUTE ON RESTART
if ( status .eq. 1 ) then
    !> REALLOCATE TO SIZE AT CHECKPOINT
    arr_c_ptr = c_loc(arr(1,1,1))
    call FTI_Realloc(0, arr_c_ptr, ierr)
    call FTI_recover(ierr)
    !> RESHAPE ARRAY
    call c_f_pointer(arr_c_ptr, arr, shape)
    call FTI_Realloc(0, arr, ierr)

! ...

end if

! ...

!> FIRST CHECKPOINT
call FTI_Checkpoint(1, 1, ierr)

! ...

!> CHANGE ARRAY DIMENSION
!> AND STORE IN SHAPE ARRAY
shape = [N21,N22,N23]
allocate(tmp(N21,N22,N23))
tmp(1:N11,1:N12,1:N13) = arr
deallocate(arr)
arr => tmp

!> TELL FTI ABOUT THE NEW DIMENSION
call FTI_Protect(0, arr, ierr)

! ...

!> SECOND CHECKPOINT
call FTI_Checkpoint(2,1, ierr)

! ...

```

For derived data-types:

```

! ...

use iso_c_binding

! ...

type polar
  real :: radius
  real :: phi
end type

type(FTI_Type)      :: FTI_Polar
integer, parameter  :: N1=128*102*25 !> 25 MB / Process
integer, parameter  :: N2=128*102*50 !> 50 MB / Process
integer, parameter  :: N11 = 128
integer, parameter  :: N12 = 102
integer, parameter  :: N13 = 25
integer, parameter  :: N21 = 128
integer, parameter  :: N22 = 102
integer, parameter  :: N23 = 50
integer, target      :: FTI_COMM_WORLD
integer              :: ierr, status

type(polar), dimension(:, :, :), pointer :: arr
type(c_ptr)          :: arr_c_ptr
type(polar), dimension(:, :, :), pointer :: tmp
integer, dimension(:), pointer          :: shape

allocate(arr(N11,N12,N13))
allocate(shape(3))

!> INITIALIZE C POINTER
arr_c_ptr = c_loc( arr( &
  lbound(arr,1), &
  lbound(arr,2), &
  lbound(arr,3)))

! ...

!> PROTECT DATA AND ITS SHAPE
call FTI_Protect(0, arr_c_ptr, size(arr), FTI_Polar, ierr)
call FTI_Protect(1, shape, ierr)

call FTI_Status(status)

!> EXECUTE ON RESTART
if ( status .eq. 1 ) then
  !> REALLOCATE TO SIZE AT CHECKPOINT
  call FTI_Realloc(0, arr_c_ptr, ierr)
  call FTI_recover(ierr)
  !> RESHAPE ARRAY
  call c_f_pointer(arr_c_ptr, arr, shape)

! ...

end if

! ...

!> FIRST CHECKPOINT
call FTI_Checkpoint(1, 1, ierr)

! ...

!> CHANGE ARRAY DIMENSION
!> AND STORE IN SHAPE ARRAY
shape = [N21,N22,N23]
allocate(tmp(N21,N22,N23))
tmp(1:N11,1:N12,1:N13) = arr
deallocate(arr)
arr => tmp

! ...

```

```

!> UPDATE C POINTER BEFORE CALL TO 'FTI_Protect'
arr_c_ptr = c_loc( arr( &
    lbound(arr,1), &
    lbound(arr,2), &
    lbound(arr,3)))

!> TELL FTI ABOUT THE NEW DIMENSION
call FTI_Protect(0, arr_c_ptr, size(arr), FTI_Polar, ierr)

! ...

!> SECOND CHECKPOINT
call FTI_Checkpoint(2,1, ierr)

! ...

```

FTI_Status

- Returns the current status of the recovery flag.

DEFINITION

```
subroutine FTI_Status ( status )
```

ARGUMENTS

Variable		What for?
integer status	OUT	Token for status flag.

OUTPUT

Value	Reason
0	No checkpoints taken yet or recovered successfully
1	At least one checkpoint is taken. If execution fails, the next start will be a restart
2	The execution is a restart from checkpoint level L4 and keep_last_checkpoint was enabled during the last execution

DESCRIPTION

FTI_Status returns the current status of the recovery flag.

EXAMPLE

```

call FTI_Status(status)

!> EXECUTE ON RESTART
if ( status .eq. 1 ) then

! ...

call FTI_recover(ierr)

! ...

end if

```

FTI_Recover

- Recovers the data of the protected variables from the checkpoint file.

DEFINITION

```
subroutine FTL_Recover ( err )
```

ARGUMENTS

Variable		What for?
integer err	OUT	Token for FTL error code.

ERROR HANDLING

Value	Reason
FTL_SCES	Success
FTL_NSCS	Failure

DESCRIPTION

FTL_Recover loads the data from the checkpoint file to the protected variables. It only recovers variables which are protected by a preceeding call to [FTL_Protect](#). If a variable changes its size during execution, the proper amount of memory has to be allocated for that variable before the call to `FTL_Recover`. FTL provides the API functions [FTL_GetStoredSize](#) and

EXAMPLE

see example of [FTL_Status](#).

FTL_Snapshot

- Invokes the recovery of protected variables on a restart.
- Writes multilevel checkpoints regarding their requested frequencies during execution.

DEFINITION

```
subroutine FTL_Snapshot ( err )
```

ARGUMENTS

Variable		What for?
integer err	OUT	Token for FTL error code.

ERROR HANDLING

Value	Reason
FTL_SCES	Successfull call (without checkpointing) or if recovery successful
FTL_NSCS	Failure of <code>FTL_Checkpoint</code>
FTL_DONE	Success of <code>FTL_Checkpoint</code>
FTL_NREC	Failure on recovery

DESCRIPTION

On a restart, `FTL_Snapshot` loads the data from the checkpoint file to the protected variables. During execution it performs checkpoints according to the checkpoint frequencies for the various safety levels. The frequencies may be set in the configuration file (see e.g.: [ckpt_L1](#)).

`FTL_Snapshot` can only take care of variables which are protected by a preceding call to [FTL_Protect](#).

EXAMPLE


```

! ...

ptriter => iter
call FTI_Protect(0, ptriter, err)
call FTI_Protect(2, g, err)
call FTI_Protect(1, h, err)

do iter = 1, ITER_TIMES

    call FTI_Snapshot(err)

    call doWork(nbProcs, rank, g, h, localerror)

! ...

enddo

if ( rank == 0 ) then
    print '("Execution finished in ",F9.0," seconds.")', MPI_Wtime() - wtime
endif

! ...

```

FTI_Finalize

- Frees the allocated memory.
- Communicates the end of the execution to dedicated threads.
- Cleans checkpoints and metadata.

DEFINITION

```
subroutine FTI_Finalize ( err )
```

ARGUMENTS

Variable		What for?
<code>integer err</code>	OUT	Token for FTI error code.

ERROR HANDLING

Value	Reason
<code>FTI_SCES</code>	For application process
<code>exit(0)</code>	For FTI process (only if head == 1)

DESCRIPTION

`FTI_Finalize` notifies the FTI processes that the execution is over, frees FTI internal data structures and it performs a clean up of the checkpoint folders at a normal execution. If the setting `keep_last_ckpt` is set, it flushes local checkpoint files (if present) to the PFS. If the setting `head` is set to 1, it will also terminate the FTI processes. It should be called before `MPI_Finalize()`.

EXAMPLE

! ...

deallocate(h)

deallocate(g)

call FTL_Finalize(err)

call MPI_Finalize(err)

! ...

[Basic]

head

The checkpointing safety levels L2, L3 and L4 produce additional overhead due to the necessary postprocessing work on the checkpoints. FTI offers the possibility to create an MPI process, called HEAD, in which this postprocessing will be accomplished. This allows it for the application processes to continue the execution immediately after the checkpointing.

Value	Meaning
0	The checkpoint postprocessing work is covered by the application processes
1	The HEAD process accomplishes the checkpoint postprocessing work (notice: In this case, the number of application processes will be (n-1)/node)

(default = 0)

node_size

Lets FTI know, how many processes will run on each node (ppn). In most cases this will be the amount of processing units within the node (e.g. 2 CPU's/node and 8 cores/CPU ! 16 processes/node).

Value	Meaning
ppn (int > 0)	Number of processing units within each node (notice: The total number of processes must be a multiple of group_size*node_size)

(default = 2)

ckpt_dir

This entry defines the path to the local hard drive on the nodes.

Value	Meaning
string	Path to the local hard drive on the nodes

(default = /scratch/username)

glbl_dir

This entry defines the path to the checkpoint folder on the PFS (L4 checkpoints).

Value	Meaning
string	Path to the checkpoint directory on the PFS

(default = /work/project)

meta_dir

This entry defines the path to the meta files directory. The directory has to be accessible from each node. It keeps files with information about the topology of the execution.

Value	Meaning
string	Path to the meta files directory

(default = /home/user/.ftt)

ckpt_L1

Here, the user sets the checkpoint frequency of L1 checkpoints when using `FTI_Snapshot()`.

Value	Meaning
L1 intv. (int >= 0)	L1 checkpointing interval in minutes
0	Disable L1 checkpointing

(default = 3)

ckpt_L2

Here, the user sets the checkpoint frequency of L2 checkpoints when using `FTI_Snapshot()`.

Value	Meaning
L2 intv. (int >= 0)	L2 checkpointing interval in minutes
0	Disable L2 checkpointing

(default = 5)

ckpt_L3

Here, the user sets the checkpoint frequency of L3 checkpoints when using `FTI_Snapshot()`.

Value	Meaning
L3 intv. (int >= 0)	L3 checkpointing interval in minutes
0	Disable L3 checkpointing

(default = 7)

ckpt_L4

Here, the user sets the checkpoint frequency of L4 checkpoints when using `FTI_Snapshot()`.

Value	Meaning
L4 intv. (int >= 0)	L4 checkpointing interval in minutes
0	Disable L4 checkpointing

(default = 11)

inline_L2

In this setting, the user chose whether the post-processing work on the L2 checkpoints is done by an FTI process or by the application process.

Value	Meaning
0	The post-processing work of the L2 checkpoints is done by an FTI process (notice: This setting is only allowed if head = 1)
1	The post-processing work of the L2 checkpoints is done by the application process

(default = 1)

inline_L3

In this setting, the user chose whether the post-processing work on the L3 checkpoints is done by an FTI process or by the application process.

Value	Meaning
0	The post-processing work of the L3 checkpoints is done by an FTI process (notice: This setting is only allowed if head = 1)
1	The post-processing work of the L3 checkpoints is done by the application process

(default = 1)

inline_L4

In this setting, the user chose whether the post-processing work on the L4 checkpoints is done by an FTI process or by the application process.

Value	Meaning
0	The post-processing work of the L4 checkpoints is done by an FTI process (notice: This setting is only allowed if head = 1)
1	The post-processing work of the L4 checkpoints is done by the application process

(default = 1)

keep_last_ckpt

This setting tells FTI whether the last checkpoint taken during the execution will be kept in the case of a successful run or not.

Value	Meaning
0	After <code>FTI_Finalize()</code> , the meta files and checkpoints will be removed. No checkpoint data will be kept on the PFS or on the local hard drives of the nodes
1	After <code>FTI_Finalize()</code> , the last checkpoint will be kept and stored on the PFS as a L4 checkpoint (notice: Additionally, the setting failure in the configuration file is set to 2. This will lead to a restart from the last checkpoint if the application is executed again)

(default = 0)

group_size

The group size entry sets, how many nodes (members) forming a group.

Value	Meaning
<div> <div>int i (2 <= i <= 32)</div> <div> <div></div> <div></div> </div> </div>	Number of nodes contained in a group (notice: The total number of processes must be a multiple of <code>group_size*node_size</code>)

(default = 4)

max_sync_intv

Sets the maximum number of iterations between synchronisations of the iteration length (used for `FTI_Snapshot()`). Internally the value will be rounded to the next lower value which is a power of 2.

Value	Meaning
<code>int i (0 <= i <= INT_MAX)</code>	maximum number of iterations between measurements of the global mean iteration time (<code>MPI_Allreduce</code> call)
<code>0</code>	Sets the value to 512, the default value for FTI

(default = 0)

ckpt_io

Sets the I/O mode.

Value	Meaning
<code>1</code>	POSIX I/O mode
<code>2</code>	MPI-IO I/O mode
<code>3</code>	SIONLib I/O mode

(default = 1)

verbosity

Sets the level of verbosity.

Value	Meaning
<code>1</code>	Debug sensitive. Beside warnings, errors and information, FTI debugging information will be printed
<code>2</code>	Information sensitive. FTI prints warnings, errors and information
<code>3</code>	FTI prints only warnings and errors
<code>4</code>	FTI prints only errors

(default = 2)

[Restart]

failure

This setting should mainly be set by FTI itself. The behaviour within FTI is the following:

- Within `FTI_Init()`, it remains on its initial value.
- After the first checkpoint is taken, it is set to 1.
- After `FTI_Finalize()` and `keep_last_ckpt = 0`, it is set to 0.
- After `FTI_Finalize()` and `keep_last_ckpt = 1`, it is set to 2.

Value	Meaning
<code>0</code>	The application starts with its initial conditions (notice: In order to force a clean start, the value may be set to 0 manually. In this case the user has to take care about removing the checkpoint data from the last execution)
<code>1</code>	FTI is searching for checkpoints and starts from the highest checkpoint level (notice: If no readable checkpoints are found, the execution stops)
<code>2</code>	FTI is searching for the last L4 checkpoint and restarts the execution from there (notice: If checkpoint is not L4 or checkpoint is not readable, the execution stops)

(default = 0)

exec_id

This setting should mainly set by FTI itself. During `FTI_Init()` the execution ID is set if the application starts for the first time (failure = 0) or the execution ID is used by FTI in order to find the checkpoint files for the case of a restart (failure = 1,2)

Value	Meaning
yyyy-mm-dd_hh-mm-ss	Execution ID (notice: If variate checkpoint data is available, the execution ID may set by the user to assign the desired starting point)

(default = NULL)

[Advanced]

The settings in this section, should **ONLY** be changed by advanced users.

block_size

FTI temporarily copies small blocks of the L2 and L3 checkpoints to send them through MPI. The size of the data blocks can be set here.

Value	Meaning
int	Size in KB of the data blocks send by FTI through MPI for the checkpoint levels L2 and L3

(default = 1024)

transfer_size

FTI transfers in chunks local checkpoint files to PFS. The size of the chunk can be set here.

Value	Meaning
int	Size in MB of the chunks send by FTI from local to PFS

(default = 16)

mpi_tag

FTI uses a certain tag for the MPI messages. This tag can be set here.

Value	Meaning
int	Tag, used for MPI messages within FTI

(default = 2612)

lustre_striping_unit

This option only impacts if `-DENABLE_LUSTRE` was added to the Cmake command. It sets the striping unit for the MPI-IO file.

Value	Meaning
int i (0 <= i <= INT_MAX)	Striping size in Bytes. The default in Lustre systems is 1MB (1048576 Bytes), FTI uses 4MB (4194304 Bytes) as the dafault value
0	Assigns the Lustre default value

(default = 4194304)

lustre_striping_factor

This option only impacts if `-DENABLE_LUSTRE` was added to the Cmake command. It sets the striping factor for the MPI-IO file.

Value	Meaning
<code>int i (0 <= i <= INT_MAX)</code>	Striping factor. The striping factor determines the number of OST's to use for striping.
<code>-1</code>	Stripe over all available OST's. This is the default in FTI.
<code>0</code>	Assigns the Lustre default value

(default = -1)

lustre_striping_offset

This option only impacts if `-DENABLE_LUSTRE` was added to the Cmake command. It sets the striping offset for the MPI-IO file.

Value	Meaning
<code>int i (0 <= i <= INT_MAX)</code>	Striping offset. The striping offset selects a particular OST to begin striping at.
<code>-1</code>	Assigns the Lustre default value

(default = -1)

local_test

FTI is building the topology of the execution, by determining the hostnames of the nodes on which each process runs. Depending on the settings for `group_size`, `node_size` and `head`, FTI assigns each particular process to a group and decides which process will be Head or Application dedicated. This is meant to be a local test. In certain situations (e.g. to run FTI on a local machine) it is necessary to disable this function.

Value	Meaning
<code>0</code>	Local test is disabled. FTI will simulate the situation set in the configuration
<code>1</code>	Local test is enabled (notice: FTI will check if the settings are correct on initialization and if necessary stop the execution)

(default = 1)

Default Configuration

```
head                = 0
node_size           = 2
ckpt_dir            = /scratch/username/
gbl_dir             = /work/project/
meta_dir            = /home/username/.fti/
ckpt_L1             = 3
ckpt_L2             = 5
ckpt_L3             = 7
ckpt_L4             = 11
inline_L2           = 1
inline_L3           = 1
inline_L4           = 1
keep_last_ckpt      = 0
group_size          = 4
max_sync_intv       = 0
ckpt_io             = 1
verbosity           = 2
failure             = 0
exec_id             = NULL
block_size          = 1024
transfer_size       = 16
mpi_tag             = 2612
lustre_stripping_unit = 4194304
lustre_stripping_factor = -1
lustre_stripping_offset = -1
local_test          = 1
```

DESCRIPTION

This configuration is made of default values (see: 5). FTI processes are not created (`head = 0`, notice: if there is no FTI processes, all post-checkpoints must be done by application processes, thus `inline_L2`, `inline_L3` and `inline_L4` are set to 1), last checkpoint won't be kept (`keep_last_ckpt = 0`), `FTI_Snapshot()` will take L1 checkpoint every 3 min, L2 - every 5 min, L3 - every 7 min and L4 - every 11 min, FTI will print errors and some few important information (`verbosity = 2`) and IO mode is set to POSIX (`ckpt_io = 1`). This is a normal launch of a job, because failure is set to 0 and `exec_id` is `NULL`. `local_test = 1` makes this a local test.

Using FTI Processes

```

head                = 1
node_size           = 2
ckpt_dir            = /scratch/username/
gbl_dir             = /work/project/
meta_dir            = /home/username/.fti/
ckpt_L1             = 3
ckpt_L2             = 5
ckpt_L3             = 7
ckpt_L4             = 11
inline_L2           = 0
inline_L3           = 0
inline_L4           = 0
keep_last_ckpt      = 0
group_size          = 4
max_sync_intv       = 0
ckpt_io             = 1
verbosity           = 2
failure             = 0
exec_id             = NULL
block_size          = 1024
transfer_size       = 16
mpi_tag             = 2612
lustre_stripping_unit = 4194304
lustre_stripping_factor = -1
lustre_stripping_offset = -1
local_test          = 1

```

DESCRIPTION

FTI processes are created (`head = 1`) and all post-checkpointing is done by them, thus `inline_L2`, `inline_L3` and `inline_L4` are set to 0. Note that it is possible to select which checkpoint levels should be post-processed by heads and which by application processes (e.g. `inline_L2 = 1`, `inline_L3 = 0`, `inline_L4 = 0`). L1 post-checkpoint is always done by application processes, because it's a local checkpoint. Be aware, when `head = 1`, and `inline_L2`, `inline_L3` and `inline_L4` are set to 1 all post-checkpoint is still made by application processes.

Using only selected ckpt level with FTI_Snapshot

```

head                = 0
node_size           = 2
ckpt_dir            = /scratch/username/
gbl_dir             = /work/project/
meta_dir            = /home/username/.fti/
ckpt_L1             = 0
ckpt_L2             = 5
ckpt_L3             = 0
ckpt_L4             = 0
inline_L2           = 1
inline_L3           = 1
inline_L4           = 1
keep_last_ckpt      = 0
group_size          = 4
max_sync_intv       = 0
ckpt_io             = 1
verbosity           = 2
failure             = 0
exec_id             = NULL
block_size          = 1024
transfer_size       = 16
mpi_tag             = 2612
lustre_stripping_unit = 4194304
lustre_stripping_factor = -1
lustre_stripping_offset = -1
local_test          = 1

```

DESCRIPTION

`FTI_Snapshot()` will take only L2 checkpoint every 5 min Notice that other configurations are also possible (e.g. take L1 ckpt every 5 min and L4 ckpt every 30 min).

Keeping last checkpoint

```
head           = 0
node_size      = 2
ckpt_dir       = /scratch/username/
gbl_dir        = /work/project/
meta_dir       = /home/username/.fti/
ckpt_L1        = 3
ckpt_L2        = 5
ckpt_L3        = 7
ckpt_L4        = 11
inline_L2      = 1
inline_L3      = 1
inline_L4      = 1
keep_last_ckpt = 1
group_size     = 4
max_sync_intv  = 0
ckpt_io        = 1
verbosity      = 2
failure        = 0
exec_id        = NULL
block_size     = 1024
transfer_size  = 16
mpi_tag        = 2612
lustre_striping_unit = 4194304
lustre_striping_factor = -1
lustre_striping_offset = -1
local_test     = 1
```

DESCRIPTION

FTI will keep last checkpoint (`Keep_last_ckpt = 1`), thus after finishing the job Failure will be set to 2.

Using different IO mode

For instance MPI-I/O:

```
head           = 0
node_size      = 2
ckpt_dir       = /scratch/username/
gbl_dir        = /work/project/
meta_dir       = /home/username/.fti/
ckpt_L1        = 3
ckpt_L2        = 5
ckpt_L3        = 7
ckpt_L4        = 11
inline_L2      = 1
inline_L3      = 1
inline_L4      = 1
keep_last_ckpt = 0
group_size     = 4
max_sync_intv  = 0
ckpt_io        = 2
verbosity      = 2
failure        = 0
exec_id        = NULL
block_size     = 1024
transfer_size  = 16
mpi_tag        = 2612
lustre_striping_unit = 4194304
lustre_striping_factor = -1
lustre_striping_offset = -1
local_test     = 1
```

DESCRIPTION

FTI IO mode is set to MPI IO (`ckpt_io = 2`). Third option is SIONlib IO mode (`ckpt_io = 3`).

Restart after a failure

```
head           = 0
node_size      = 2
ckpt_dir       = /scratch/username/
glibl_dir      = /work/project/
meta_dir       = /home/username/.fti/
ckpt_L1        = 3
ckpt_L2        = 5
ckpt_L3        = 7
ckpt_L4        = 11
inline_L2      = 1
inline_L3      = 1
inline_L4      = 1
keep_last_ckpt = 0
group_size     = 4
max_sync_intv  = 0
ckpt_io        = 1
verbosity      = 2
failure        = 1
exec_id        = 2017-07-26_13-22-11
block_size     = 1024
transfer_size  = 16
mpi_tag        = 2612
lustre_striping_unit = 4194304
lustre_striping_factor = -1
lustre_striping_offset = -1
local_test     = 1
```

DESCRIPTION

This config tells FTI that this job is a restart after a failure (`failure` set to 1 and `exec_id` is some date in a format YYYY-MM-DD_HH-mm-ss , where YYYY - year, MM - month, DD - day, HH - hours, mm - minutes, ss - seconds). When recovery is not possible, FTI will abort the job (when using `FTI_Snapshot()`) and/or signal failed recovery by `FTI_Status()` .

Using FTI_Snapshot

```
#include <stdlib.h>
#include <fti.h>

int main(int argc, char** argv){
    MPI_Init(&argc, &argv);
    char* path = "config.fti"; //config file path
    FTI_Init(path, MPI_COMM_WORLD);
    int world_rank, world_size; //FTI_COMM rank & size
    MPI_Comm_rank(FTI_COMM_WORLD, &world_rank);
    MPI_Comm_size(FTI_COMM_WORLD, &world_size);

    int *array = malloc(sizeof(int) * world_size);
    int number = world_rank;
    int i = 0;
    //adding variables to protect
    FTI_Protect(1, &i, 1, FTI_INTG);
    FTI_Protect(2, &number, 1, FTI_INTG);
    for (; i < 100; i++) {
        FTI_Snapshot();
        MPI_Allgather(&number, 1, MPI_INT, array,
            1, MPI_INT, FTI_COMM_WORLD);
        number += 1;
    }
    free(array);
    FTI_Finalize();
    MPI_Finalize();
    return 0;
}
```

DESCRIPTION

`FTI_Snapshot()` makes a checkpoint by given time and also recovers data after a failure, thus makes the code shorter. Checkpoints intervals can be set in configuration file (see: [ckpt_L1](#) - [ckpt_L4](#)).

Using FTI_Checkpoint

```

#include <stdlib.h>
#include <fti.h>
#define ITER_CHECK 10

int main(int argc, char** argv){
    MPI_Init(&argc, &argv);
    char* path = "config.fti"; //config file path
    FTI_Init(path, MPI_COMM_WORLD);
    int world_rank, world_size; //FTI_COMM rank & size
    MPI_Comm_rank(FTI_COMM_WORLD, &world_rank);
    MPI_Comm_size(FTI_COMM_WORLD, &world_size);

    int *array = malloc(sizeof(int) * world_size);
    int number = world_rank;
    int i = 0;
    //adding variables to protect
    FTI_Protect(1, &i, 1, FTI_INTG);
    FTI_Protect(2, &number, 1, FTI_INTG);
    if (FTI_Status() != 0) {
        FTI_Recover();
    }
    for (; i < 100; i++) {
        if (i % ITER_CHECK == 0) {
            FTI_Checkpoint(i / ITER_CHECK + 1, 2);
        }
        MPI_Allgather(&number, 1, MPI_INT, array,
            1, MPI_INT, FTI_COMM_WORLD);
        number += 1;
    }
    free(array);
    FTI_Finalize();
    MPI_Finalize();
    return 0;
}

```

DESCRIPTION

FTI_Checkpoint() allows to checkpoint at precise application intervals. Note that when using FTI_Checkpoint(), ckpt_L1, ckpt_L2, ckpt_L3 and ckpt_L4 are not taken into account.

Using FTI_Realloc with Fortran and Intrinsic Types

```

program test_fti_realloc
  use fti
  use iso_c_binding
  implicit none
  include 'mpif.h'

  integer, parameter      :: dp=kind(1.0d0)
  integer, parameter      :: N1=128*1024*25 !> 25 MB / Process
  integer, parameter      :: N2=128*1024*50 !> 50 MB / Process
  integer, parameter      :: N11 = 128
  integer, parameter      :: N12 = 1024
  integer, parameter      :: N13 = 25
  integer, parameter      :: N21 = 128
  integer, parameter      :: N22 = 1024
  integer, parameter      :: N23 = 50
  integer, target         :: FTI_COMM_WORLD
  integer                 :: ierr, status

  real(dp), dimension(:, :, :), pointer :: arr
  type(c_ptr)            :: arr_c_ptr
  real(dp), dimension(:, :, :), pointer :: tmp
  integer(4), dimension(:), pointer :: shape

  allocate(arr(N11,N12,N13))
  allocate(shape(3))

  !> INITIALIZE MPI AND FTI
  call MPI_Init(ierr)
  FTI_COMM_WORLD = MPI_COMM_WORLD
  call FTI_Init('config.fti', FTI_COMM_WORLD, ierr)

  !> PROTECT DATA AND ITS SHAPE
  call FTI_Protect(0, arr, ierr)
  call FTI_Protect(1, shape, ierr)

  call FTI_Status(status)

  !> EXECUTE ON RESTART
  if ( status .eq. 1 ) then
    !> REALLOCATE TO SIZE AT CHECKPOINT
    arr_c_ptr = c_loc(arr(1,1,1))
    call FTI_Realloc(0, arr_c_ptr, ierr)
    call FTI_recover(ierr)
    !> RESHAPE ARRAY
    call c_f_pointer(arr_c_ptr, arr, shape)
    call FTI_Finalize(ierr)
    call MPI_Finalize(ierr)
    STOP
  end if

  !> FIRST CHECKPOINT
  call FTI_Checkpoint(1, 1, ierr)

  !> CHANGE ARRAY DIMENSION
  !> AND STORE IN SHAPE ARRAY
  shape = [N21,N22,N23]
  allocate(tmp(N21,N22,N23))
  tmp(1:N11,1:N12,1:N13) = arr
  deallocate(arr)
  arr => tmp

  !> TELL FTI ABOUT THE NEW DIMENSION
  call FTI_Protect(0, arr, ierr)

  !> SECOND CHECKPOINT
  call FTI_Checkpoint(2, 1, ierr)

  !> SIMULATE CRASH
  call MPI_Abort(MPI_COMM_WORLD,-1,ierr)
end program

```

Using FTI_Realloc with Fortran and Derived Types

```

program test_fti_realloc
  use fti
  use iso_c_binding
  implicit none
  include 'mpif.h'

  !> DEFINE DERIVED TYPE
  type :: polar
    real :: radius
    real :: phi
  end type

  integer, parameter :: dp=kind(1.0d0)
  integer, parameter :: N1=128*1024*25 !> 25 MB / Process
  integer, parameter :: N2=128*1024*50 !> 50 MB / Process
  integer, parameter :: N11 = 128
  integer, parameter :: N12 = 1024
  integer, parameter :: N13 = 25
  integer, parameter :: N21 = 128
  integer, parameter :: N22 = 1024
  integer, parameter :: N23 = 50
  integer, target :: FTI_COMM_WORLD
  integer :: ierr, status
  type(FTI_type) :: FTI_Polar

  type(c_ptr) :: cPtr
  type(polar), dimension(:, :, :), pointer :: arr
  type(polar), dimension(:, :, :), pointer :: tmp
  integer(4), dimension(:), pointer :: shape

  !> INITIALIZE FTI TYPE 'FTI_POLAR'
  call FTI_InitType(FTI_Polar, 2*4, ierr)

  allocate(arr(N11,N12,N13))
  allocate(shape(3))

  !> INITIALIZE MPI AND FTI
  call MPI_Init(ierr)
  FTI_COMM_WORLD = MPI_COMM_WORLD
  call FTI_Init('config.fti', FTI_COMM_WORLD, ierr)

  !> PROTECT DATA AND ITS SHAPE
  call FTI_Protect(0, c_loc(arr), size(arr), FTI_Polar, ierr)
  call FTI_Protect(1, shape, ierr)

  call FTI_Status(status)

  !> EXECUTE ON RESTART
  if ( status .eq. 1 ) then
    !> REALLOCATE TO DIMENSION AT LAST CHECKPOINT
    cPtr = c_loc(arr)
    call FTI_Realloc(0, cPtr, ierr) !> PASS DATA AS C-POINTER
    call FTI_recover(ierr)
    call c_f_pointer(cPtr, arr, shape) !> CAST BACK TO F-POINTER
    call FTI_Finalize(ierr)
    call MPI_Finalize(ierr)
    STOP
  end if

  !> FIRST CHECKPOINT
  call FTI_Checkpoint(1, 1, ierr)

  !> CHANGE ARRAY DIMENSION
  !> AND STORE IN SHAPE ARRAY
  shape = [N21,N22,N23]
  allocate(tmp(N21,N22,N23))
  tmp(1:N11,1:N12,1:N13) = arr
  deallocate(arr)
  arr => tmp

  !> TELL FTI ABOUT THE NEW DIMENSION
  call FTI_Protect(0, c_loc(arr), size(arr), FTI_Polar, ierr)

```



```
!> SECOND CHECKPOINT
```

```
call FTI_Checkpoint(2,1, ierr)
```

```
!> SIMULATE CRASH
```

```
call MPI_Abort(MPI_COMM_WORLD,-1,ierr)
```

```
end program
```

List of applications integrated with FTI

CoMD

Classical molecular dynamics proxy application.

<https://github.com/exmatex/CoMD>

File changes

Integrating FTI in CoMD took only addition of ~30 lines of code in 2 files. All occurrences of `MPI_COMM_WORLD` changed to `FTI_COMM_WORLD` except `FTI_Init("config.fti", MPI_COMM_WORLD);`

```
File: src-mpi/CoMD.c
102:  int i = 1;
103:  FTI_Protect(i++, sim->boxes->nAtoms, sim->boxes->nTotalBoxes, FTI_INTG);
104:
105:  FTIT_type RealTInfo;
106:  FTI_InitType(&RealTInfo, sizeof(real_t));
107:  FTIT_type Real3Info;
108:  FTI_InitType(&Real3Info, sizeof(real3));
109:  int maxTotalAtoms = MAXATOMS * (sim->boxes->nTotalBoxes);
110:
111:  FTI_Protect(i++, sim->atoms->gid, maxTotalAtoms, FTI_INTG);
112:  FTI_Protect(i++, sim->atoms->iSpecies, maxTotalAtoms, FTI_INTG);
113:  FTI_Protect(i++, sim->atoms->r, maxTotalAtoms, Real3Info);
114:  FTI_Protect(i++, sim->atoms->p, maxTotalAtoms, Real3Info);
115:  FTI_Protect(i++, sim->atoms->f, maxTotalAtoms, Real3Info);
116:
117:  int iStep = 0;
118:  FTI_Protect(i++, &iStep, 1, FTI_INTG);
119:
120:  if (FTI_Status() != 0) {
121:      int res = FTI_Recover();
122:      if (res != 0) {
123:          printf("\tRecovery failed! FTI_Recover returned %d.\n", res);
124:      }
125:  }
126:  ---
139:  profileStart(loopTimer);
140:  for (; iStep < nSteps; )
141:  {
142:      startTimer(commReduceTimer);
143:      sumAtoms(sim);
144:      stopTimer(commReduceTimer);
145:
146:      printThings(sim, iStep, getElapsedTime(timestepTimer));
147:
148:      startTimer(timestepTimer);
149:      timestep(sim, printRate, sim->dt);
150:      stopTimer(timestepTimer);
151:
152:      iStep += printRate;
153:      int res = FTI_Checkpoint(iStep, 1);
154:      if (res != FTI_DONE) {
155:          printf("\tCheckpoint failed! FTI_Checkpoint returned %d.\n", res);
156:      }
157:  }
158:  profileStop(loopTimer);
```

```

File: src-mpi/parallel.c
64: void initParallel(int* argc, char*** argv)
65: {
66: #ifdef DO_MPI
67:   MPI_Init(argc, argv);
68:   FTL_Init("config.fti", MPI_COMM_WORLD);
69:   MPI_Comm_rank(FTL_COMM_WORLD, &myRank);
70:   MPI_Comm_size(FTL_COMM_WORLD, &nRanks);
71: #endif
72: }
73:
74: void destroyParallel()
75: {
76: #ifdef DO_MPI
77:   FTL_Finalize();
78:   MPI_Finalize();
79: #endif
80: }

```

Results

Log of run without FTI integrated.

```

=====
Poznan Supercomputing and Networking Center
eagle.man.poznan.pl
=====

Start of calculations [pon, 9 paź 2017, 12:57:49 CEST]

Support:      support-hpc@man.poznan.pl

Mon Oct 9 12:57:53 2017: Starting Initialization
Mini-Application Name   : CoMD-mpi
Mini-Application Version : 1.1
Platform:
  hostname: e0026
  kernel name: 'Linux'
  kernel release: '3.10.105-1.el6.elrepo.x86_64'
  processor: 'x86_64'
Build:
  CC: '/opt/exp_soft/local/generic/openmpi/1.10.2-1_gcc482/bin/mpicc'
  compiler version: 'gcc (GCC) 4.8.2 20140120 (Red Hat 4.8.2-14)'
  CFLAGS: '-std=c99 -DDOUBLE -DDO_MPI -g -O5 -I/home/users/ksiero1/fti/include/ '
  LDFLAGS: '-L/home/users/ksiero1/fti/lib/ -lm -lcrypto'
  using MPI: true
  Threading: none
  Double Precision: true
Run Date/Time: 2017-10-09, 12:57:53
Command Line Parameters:
doeam: 0
potDir: pots
potName: Cu_u6.eam
potType: funcfl
nx: 800
ny: 800
nz: 800
xproc: 8
yproc: 8
zproc: 8
Lattice constant: -1 Angstroms
nSteps: 100
printRate: 10
Time step: 1 fs
Initial Temperature: 600 K
Initial Delta: 0 Angstroms
Simulation data:
  Total atoms      : 2048000000
  Min global bounds : [ 0.00000000000, 0.00000000000, 0.00000000000 ]
  Max global bounds : [ 2892.00000000000, 2892.00000000000, 2892.00000000000 ]
Decomposition data:
  Processors      : 8. 8. 8

```

Local boxes : 62, 62, 62 = 238328
Box size : [5.8306451613, 5.8306451613, 5.8306451613]
Box factor : [1.0074548875, 1.0074548875, 1.0074548875]
Max Link Cell Occupancy: 32 of 64

Potential data:

Potential type : Lennard-Jones
Species name : Cu
Atomic number : 29
Mass : 63.55 amu
Lattice Type : FCC
Lattice spacing : 3.615 Angstroms
Cutoff : 5.7875 Angstroms
Epsilon : 0.167 eV
Sigma : 2.315 Angstroms

Memory data:

Intrinsic atom footprint = 88 B/atom
Total atom footprint = -157.000 MB (335.69 MB/node)
Link cell atom footprint = 1280.082 MB/node
Link cell atom footprint = 1408.000 MB/node (including halo cell data)

Initial energy : -1.166063303598, atom count : 2048000000

Mon Oct 9 12:58:06 2017: Initialization Finished

Mon Oct 9 12:58:06 2017: Starting simulation

#	Performance						
# Loop	Time(fs)	Total Energy	Potential Energy	Kinetic Energy	Temperature (us/atom)	# Atoms	
0	0.00	-1.166063303598	-1.243619295198	0.077555991600	600.0000	0.0000 2048000000	
10	10.00	-1.166059649733	-1.233151964368	0.067092314635	519.0494	2.3384 2048000000	
20	20.00	-1.166048425247	-1.208164731096	0.042116305849	325.8263	2.4122 2048000000	
30	30.00	-1.166037572103	-1.186566075400	0.020528503297	158.8156	2.4182 2048000000	
40	40.00	-1.166042088520	-1.183621872290	0.017579783770	136.0033	2.4197 2048000000	
50	50.00	-1.166051685771	-1.193725983586	0.027674297815	214.0979	2.4213 2048000000	
60	60.00	-1.166054644001	-1.202677534791	0.036622890790	283.3274	2.4201 2048000000	
70	70.00	-1.166052134038	-1.204922829363	0.038870695326	300.7172	2.4207 2048000000	
80	80.00	-1.166048793793	-1.203643980438	0.037595186645	290.8494	2.4198 2048000000	
90	90.00	-1.166048002607	-1.203830919192	0.037782916585	292.3017	2.4193 2048000000	
100	100.00	-1.166049790544	-1.206871500823	0.040821710279	315.8109	2.4176 2048000000	

Mon Oct 9 13:14:11 2017: Ending simulation

Simulation Validation:

Initial energy : -1.166063303598
Final energy : -1.166049790544
eFinal/eInitial : 0.999988
Final atom count : 2048000000, no atoms lost

Timings for Rank 0

Timer	# Calls	Avg/Call (s)	Total (s)	% Loop
total	1	977.2662	977.2662	101.34
loop	1	964.3040	964.3040	100.00
timestep	10	96.4285	964.2854	100.00
position	100	0.1001	10.0087	1.04
velocity	200	0.1025	20.5055	2.13
redistribute	101	1.2731	128.5869	13.33
atomHalo	101	0.9613	97.0902	10.07
force	101	7.8922	797.1134	82.66
commHalo	303	0.3041	92.1548	9.56
commReduce	39	0.4388	17.1143	1.77

Timing Statistics Across 512 Ranks:

Timer	Rank: Min(s)	Rank: Max(s)	Avg(s)	Stdev(s)
total	51: 977.2630	140: 977.2672	977.2650	0.0012
loop	67: 964.3023	104: 964.3043	964.3033	0.0008
timestep	51: 964.2738	463: 964.2999	964.2841	0.0061
position	49: 4.5466	373: 16.3177	11.2964	3.3782
velocity	3: 7.8438	329: 29.0049	20.1119	6.3239
redistribute	51: 53.6754	481: 168.3860	127.0497	17.6334
atomHalo	51: 24.1044	481: 142.1157	94.3223	21.2231
force	323: 775.9471	51: 905.1509	795.6917	9.6434
commHalo	51: 19.3264	481: 137.4904	89.3663	21.3808
commReduce	51: 2.8272	339: 27.7147	19.1658	4.1010

Average atom update rate: 2.41 us/atom/task

Average all atom update rate: 0.00 us/atom

Average atom rate: 212.39 atoms/us

Mon Oct 9 13:14:11 2017: CoMD Ending

End of calculations [pon, 9 paź 2017, 13:14:11 CEST].

Log of run with FTI integrated.

```
=====
Poznan Supercomputing and Networking Center
eagle.man.poznan.pl
=====

Start of calculations [pon, 9 paź 2017, 12:14:02 CEST]

Support: support-hpc@man.poznan.pl

Mon Oct 9 12:14:08 2017: Starting Initialization
Mini-Application Name : CoMD-mpi
Mini-Application Version : 1.1
Platform:
  hostname: e0026
  kernel name: 'Linux'
  kernel release: '3.10.105-1.el6.elrepo.x86_64'
  processor: 'x86_64'
Build:
  CC: '/opt/exp_soft/local/generic/openmpi/1.10.2-1_gcc482/bin/mpicc'
  compiler version: 'gcc (GCC) 4.8.2 20140120 (Red Hat 4.8.2-14)'
  CFLAGS: '-std=c99 -DDOUBLE -DDO_MPI -g -O5 -I/home/users/ksiero1/fti/include/ '
  LDFLAGS: '-L/home/users/ksiero1/fti/lib/ -lm -lcrypto'
  using MPI: true
  Threading: none
  Double Precision: true
Run Date/Time: 2017-10-09, 12:14:08
Command Line Parameters:
  doeam: 0
  potDir: pots
  potName: Cu_u6.eam
  potType: funcfl
  nx: 800
  ny: 800
  nz: 800
  xproc: 8
  yproc: 8
  zproc: 8
  Lattice constant: -1 Angstroms
  nSteps: 100
  printRate: 10
  Time step: 1 fs
  Initial Temperature: 600 K
  Initial Delta: 0 Angstroms
Simulation data:
  Total atoms : 2048000000
  Min global bounds : [ 0.0000000000, 0.0000000000, 0.0000000000 ]
  Max global bounds : [ 2892.0000000000, 2892.0000000000, 2892.0000000000 ]
Decomposition data:
  Processors : 8, 8, 8
  Local boxes : 62, 62, 62 = 238328
  Box size : [ 5.8306451613, 5.8306451613, 5.8306451613 ]
  Box factor : [ 1.0074548875, 1.0074548875, 1.0074548875 ]
  Max Link Cell Occupancy: 32 of 64
Potential data:
  Potential type : Lennard-Jones
  Species name : Cu
  Atomic number : 29
  Mass : 63.55 amu
  Lattice Type : FCC
  Lattice spacing : 3.615 Angstroms
  Cutoff : 5.7875 Angstroms
  Epsilon : 0.167 eV
  Sigma : 2.315 Angstroms
Memory data:
```

Intrinsic atom footprint = 88 B/atom
 Total atom footprint = -157.000 MB (335.69 MB/node)
 Link cell atom footprint = 1280.082 MB/node
 Link cell atom footprint = 1408.000 MB/node (including halo cell data)
 Initial energy : -1.166063303598, atom count : 2048000000
 Mon Oct 9 12:14:21 2017: Initialization Finished
 Mon Oct 9 12:14:21 2017: Starting simulation

#		Performance					
#	Loop	Time(fs)	Total Energy	Potential Energy	Kinetic Energy	Temperature (us/atom)	# Atoms
0	0.00	-1.166063303598	-1.243619295198	0.077555991600	600.0000	0.0000	2048000000
10	10.00	-1.166059649733	-1.233151964368	0.067092314635	519.0494	2.4020	2048000000
20	20.00	-1.166048425247	-1.208164731096	0.042116305849	325.8263	2.4509	2048000000
30	30.00	-1.166037572103	-1.186566075400	0.020528503297	158.8156	2.5215	2048000000
40	40.00	-1.166042088520	-1.183621872290	0.017579783770	136.0033	2.2746	2048000000
50	50.00	-1.166051685771	-1.193725983586	0.027674297815	214.0979	2.2883	2048000000
60	60.00	-1.166054644001	-1.202677534791	0.036622890790	283.3274	2.3185	2048000000

mpirun has exited due to process rank 416 with PID 0 on
 node e0769 exiting improperly. There are three reasons this could occur:

1. this process did not call "init" before exiting, but others in the job did. This can cause a job to hang indefinitely while it waits for all processes to call "init". By rule, if one process calls "init", then ALL processes must call "init" prior to termination.
2. this process called "init", but exited without calling "finalize". By rule, all processes that call "init" MUST call "finalize" prior to exiting or it will be considered an "abnormal termination"
3. this process called "MPI_Abort" or "orte_abort" and the mca parameter orte_create_session_dirs is set to false. In this case, the run-time cannot detect that the abort call was an abnormal termination. Hence, the only error message you will receive is this one.

This may have caused other processes in the application to be terminated by signals sent by mpirun (as reported here).

You can avoid this message by specifying -quiet on the mpirun command line.

Mon Oct 9 12:30:09 2017: Starting Initialization

Mini-Application Name : CoMD-mpi

Mini-Application Version : 1.1

Platform:

hostname: e0026

kernel name: 'Linux'

kernel release: '3.10.105-1.el6.elrepo.x86_64'

processor: 'x86_64'

Build:

CC: '/opt/exp_soft/local/generic/openmpi/1.10.2-1_gcc482/bin/mpicc'

compiler version: 'gcc (GCC) 4.8.2 20140120 (Red Hat 4.8.2-14)'

CFLAGS: '-std=c99 -DDOUBLE -DDO_MPI -g -O5 -I/home/users/ksiero1/fti/include/ '

LDFLAGS: '-L/home/users/ksiero1/fti/lib/ -lm -lcrypto'

using MPI: true

Threading: none

Double Precision: true

Run Date/Time: 2017-10-09, 12:30:09

Command Line Parameters:

doeam: 0

potDir: pots

potName: Cu_u6.eam

potType: funcfl

nx: 800

ny: 800

nz: 800

xproc: 8

yproc: 8

zproc: 8

Lattice constant: -1 Angstroms

nSteps: 100

printRate: 10

Time step: 1 fs

Initial Temperature: 600 K

Initial Delta: 0 Angstroms

Simulation data:

Total atoms : 2048000000

Min global bounds : [0.0000000000, 0.0000000000, 0.0000000000]

Max global bounds : [2892.0000000000, 2892.0000000000, 2892.0000000000]

Decomposition data:

Processors : 8, 8, 8

Local boxes : 63, 63, 63, 228228

Local boxes : 02, 02, 02 = 200320

Box size : [5.8306451613, 5.8306451613, 5.8306451613]

Box factor : [1.0074548875, 1.0074548875, 1.0074548875]

Max Link Cell Occupancy: 32 of 64

Potential data:

Potential type : Lennard-Jones

Species name : Cu

Atomic number : 29

Mass : 63.55 amu

Lattice Type : FCC

Lattice spacing : 3.615 Angstroms

Cutoff : 5.7875 Angstroms

Epsilon : 0.167 eV

Sigma : 2.315 Angstroms

Memory data:

Intrinsic atom footprint = 88 B/atom

Total atom footprint = -157.000 MB (335.69 MB/node)

Link cell atom footprint = 1280.082 MB/node

Link cell atom footprint = 1408.000 MB/node (including halo cell data)

Initial energy : -1.166063303598, atom count : 2048000000

Mon Oct 9 12:30:22 2017: Initialization Finished

Mon Oct 9 12:30:22 2017: Starting simulation

#					Performance		
#	Loop	Time(fs)	Total Energy	Potential Energy	Kinetic Energy	Temperature (us/atom)	# Atoms
70	70.00	-1.166063303598	-1.243619295198	0.077555991600	600.0000	0.0000	2048000000
80	80.00	-1.166048793793	-1.203643980438	0.037595186645	290.8494	2.2586	2048000000
90	90.00	-1.166048002607	-1.203830919192	0.037782916585	292.3017	2.3377	2048000000
100	100.00	-1.166049790544	-1.206871500823	0.040821710279	315.8109	2.3146	2048000000

Mon Oct 9 12:36:51 2017: Ending simulation

Simulation Validation:

Initial energy : -1.166063303598

Final energy : -1.166049790544

eFinal/eInitial : 0.999988

Final atom count : 2048000000, no atoms lost

Timings for Rank 0

Timer	# Calls	Avg/Call (s)	Total (s)	% Loop
total	1	401.7819	401.7819	103.47
loop	1	388.3197	388.3197	100.00
timestep	3	92.1447	276.4340	71.19
position	30	0.1194	3.5818	0.92
velocity	60	0.1042	6.2535	1.61
redistribute	31	0.8148	25.2584	6.50
atomHalo	31	0.4576	14.1847	3.65
force	31	8.0059	248.1816	63.91
commHalo	93	0.1349	12.5416	3.23
commReduce	18	0.1819	3.2744	0.84

Timing Statistics Across 512 Ranks:

Timer	Rank: Min(s)	Rank: Max(s)	Avg(s)	Stdev(s)
total	37: 401.7036	42: 401.9202	401.7762	0.0357
loop	34: 388.3196	370: 388.4222	388.3436	0.0197
timestep	79: 276.2655	235: 276.5081	276.4445	0.0336
position	147: 1.3705	221: 6.0085	3.6971	0.9796
velocity	147: 2.3765	206: 9.8759	6.5194	1.7920
redistribute	206: 18.4708	417: 37.1847	25.1751	4.2695
atomHalo	415: 5.8228	417: 29.1367	13.9548	5.2873
force	481: 241.1729	10: 261.1229	247.7969	2.5561
commHalo	415: 4.1255	417: 27.6639	12.3329	5.3435
commReduce	415: 1.4558	193: 6.4059	3.3989	0.9957

Average atom update rate: 2.30 us/atom/task

Average all atom update rate: 0.00 us/atom

Average atom rate: 222.25 atoms/us

Mon Oct 9 12:36:52 2017: CoMD Ending

End of calculations [pon, 9 paź 2017, 12:36:53 CEST].

CoSP2

Linear algebra algorithms and workloads for a quantum molecular dynamics (QMD) electronic structure code.

<https://github.com/exmatex/CoSP2>

File changes

Integrating FTL in CoSP2 took only addition of ~30 lines of code in 2 files. All occurrences of `MPI_COMM_WORLD` changed to `FTL_COMM_WORLD` except `FTL_Init("config.fti", MPI_COMM_WORLD);`

```
File: src-mpi/sp2Loop.c
Function: sp2Loop()
56:  FTIT_type RealTInfo;
57:  FTL_InitType(&RealTInfo, sizeof(real_t));
58:  int i = 1;
59:  FTL_Protect(i++, &iter, 1, FTL_INTG);
60:  FTL_Protect(i++, xmatrix->ia, xmatrix->hsize, FTL_INTG);
61:  FTL_Protect(i++, xmatrix->jjcontig, xmatrix->hsize * xmatrix->msize, FTL_INTG);
62:  FTL_Protect(i++, xmatrix->valcontig, xmatrix->hsize * xmatrix->msize, RealTInfo);
63:
64:  if (FTL_Status() != 0) {
65:      int res = FTL_Recover();
66:      if (res != 0) {
67:          printf("\tRecovery failed! FTL_Recover returned %d.\n", res);
68:      }
69:  }
70:
...
153:  if (iter % 10 == 0) {
154:      int res = FTL_Checkpoint(iter, 1);
155:      if (res != FTL_DONE) {
156:          printf("\tCheckpoint failed! FTL_Checkpoint returned %d.\n", res);
157:      }
158:  }
```

```
File: src-mpi/parallel.c
70: void initParallel(int* argc, char*** argv)
71: {
72: #ifdef DO_MPI
73:  MPI_Init(argc, argv);
74:  FTL_Init("config.fti", MPI_COMM_WORLD);
75:  MPI_Comm_rank(MPI_COMM_WORLD, &myRank);
76:  MPI_Comm_size(MPI_COMM_WORLD, &nRanks);
77:
78:  requestList = (MPI_Request*) malloc(nRanks*sizeof(MPI_Request));
79:  rUsed = (int*) malloc(nRanks*sizeof(int));
80:  for (int i = 0; i < nRanks; i++) { rUsed[i] = 0; }
81: #endif
82: }
83:
84: void destroyParallel()
85: {
86: #ifdef DO_MPI
87:  free(requestList);
88:  FTL_Finalize();
89:  MPI_Finalize();
90: #endif
91: }
```

Results

Log of run without FTL integrated.

```
=====
Poznan Supercomputing and Networking Center

eagle.man.poznan.pl
=====
```


Support: support-hpc@man.poznan.pl

CoSP2: SP2 Loop

Parameters:

mnpars = 80 hDim = 98304 debug = 1

hmatName =

eps = 1e-05 hEps = 1e-16

idemTol = 1e-14

hDim = 98304 M = 80

Adjusted M = 96

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 3 local row min = 18432 row max = 24576 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 1 local row min = 6144 row max = 12288 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 11 local row min = 67584 row max = 73728 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 13 local row min = 79872 row max = 86016 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 8 local row min = 49152 row max = 55296 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 7 local row min = 43008 row max = 49152 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 6 local row min = 36864 row max = 43008 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 4 local row min = 24576 row max = 30720 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 9 local row min = 55296 row max = 61440 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 10 local row min = 61440 row max = 67584 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 14 local row min = 86016 row max = 92160 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 12 local row min = 73728 row max = 79872 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 2 local row min = 12288 row max = 18432 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

total procs = 16 total rows = 98304 total cols = 96

global row min = 0 row max = 98304 row extent = 98304

rank = 0 local row min = 0 row max = 6144 row extent = 6144

Sparsity:

Initial sparsity = 672042, fraction = 6.258879e-04, Avg per row = 6.836365

Max per row = 7

l = 4, count = 2, fraction = 0.000020

l = 5, count = 621, fraction = 0.006317

l = 6, count = 14838, fraction = 0.150940

l = 7, count = 82843, fraction = 0.842723

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 15 local row min = 92160 row max = 98304 row extent = 6144

Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 5 local row min = 30720 row max = 36864 row extent = 6144

Gershgorin:

New eMax, eMin = 1.745500e+00, -7.356212e-01

bufferSize = 9437184

Initial sparsity normalized = 672042, fraction = 6.258879e-04, avg = 6.83636, max = 7

SP2Loop:

iter = 0 trX = 4.935743e+04 trX2 = 2.720037e+04

iter = 1 trX = 2.720037e+04 trX2 = 9.994787e+03

iter = 2 trX = 4.440595e+04 trX2 = 2.485384e+04

iter = 3 trX = 6.395806e+04 trX2 = 4.735425e+04

iter = 4 trX = 4.735425e+04 trX2 = 3.149323e+04

iter = 5 trX = 6.321528e+04 trX2 = 5.026180e+04

iter = 6 trX = 5.026180e+04 trX2 = 3.881328e+04

iter = 7 trX = 3.881328e+04 trX2 = 2.922713e+04

iter = 8 trX = 4.839943e+04 trX2 = 4.062611e+04

iter = 9 trX = 5.617275e+04 trX2 = 4.981154e+04

iter = 10 trX = 4.981154e+04 trX2 = 4.464542e+04

iter = 11 trX = 4.464542e+04 trX2 = 4.032639e+04

iter = 12 trX = 4.032639e+04 trX2 = 3.554445e+04

```

iter = 12 trX = 4.896445e+04 trX2 = 4.554145e+04
iter = 13 trX = 5.238745e+04 trX2 = 4.956883e+04
iter = 14 trX = 4.956883e+04 trX2 = 4.731790e+04
iter = 15 trX = 4.731790e+04 trX2 = 4.544718e+04
iter = 16 trX = 4.918861e+04 trX2 = 4.771064e+04
iter = 17 trX = 4.771064e+04 trX2 = 4.649398e+04
iter = 18 trX = 4.892731e+04 trX2 = 4.795556e+04
iter = 19 trX = 4.989906e+04 trX2 = 4.910173e+04
iter = 20 trX = 4.910173e+04 trX2 = 4.855031e+04
iter = 21 trX = 4.965316e+04 trX2 = 5.060054e+04
iter = 22 trX = 4.870578e+04 trX2 = -9.750371e+05
iter = 23 trX = 1.072449e+06 trX2 = -5.136388e+12
iter = 24 trX = -5.136388e+12 trX2 = 7.295617e+24

```

Results:

X2 Sparsity CCN = 2906510, fraction = 2.706898e-03 avg = 29.5665, max = 89
D Sparsity AAN = 2906464, fraction = 2.706856e-03 avg = 29.5661, max = 89
Number of iterations = 25

Counters for Rank 0

Counter	Calls	Avg/Call(MB)	Total(MB)
reduce	29	0.0000	0.0004
send	39	2.2910	89.3504
recv	39	2.2772	88.8095

Counter Statistics Across 16 Ranks:

Counter	Rank: Min(MB)	Rank: Max(MB)	Avg(MB)	Stdev(MB)
reduce	0: 0.0004	0: 0.0004	0.0004	0.0000
send	15: 87.4100	7: 138.5495	129.9097	15.7564
recv	15: 88.4093	6: 137.0340	129.9097	15.6236

Timings for Rank 0

Timer	# Calls	Avg/Call (s)	Total (s)	% Loop
total	1	3.4711	3.4711	100.00
loop	1	3.4711	3.4711	100.00
pre	1	0.5444	0.5444	15.68
sp2Loop	1	2.7193	2.7193	78.34
norm	1	0.0417	0.0417	1.20
x2	25	0.0473	1.1820	34.05
xadd	13	0.0454	0.5899	16.99
xset	12	0.0383	0.4591	13.23
exchange	50	0.0032	0.1576	4.54
reduceComm	29	0.0070	0.2034	5.86

Timing Statistics Across 16 Ranks:

Timer	Rank: Min(s)	Rank: Max(s)	Avg(s)	Stdev(s)
total	1: 3.4591	15: 3.5566	3.5160	0.0296
loop	1: 3.4591	15: 3.5566	3.5160	0.0296
pre	3: 0.4203	5: 0.5927	0.5180	0.0440
sp2Loop	15: 2.7191	12: 2.7256	2.7229	0.0019
norm	3: 0.0082	7: 0.0450	0.0376	0.0112
x2	1: 0.2678	15: 1.1916	1.0701	0.3027
xadd	1: 0.0548	0: 0.5899	0.5167	0.1744
xset	1: 0.0408	15: 0.4638	0.4071	0.1383
exchange	0: 0.1576	1: 1.1532	0.3589	0.2991
reduceComm	5: 0.0513	3: 1.4170	0.3006	0.4217

End of calculations [pon, 16 paź 2017, 12:17:13 CEST].

Log of run with FTI integrated.

Support: support-hpc@man.poznan.pl

CoSP2: SP2 Loop

Parameters:

msparse = 80 hDim = 98304 debug = 1
hmatName =
eps = 1e-05 hEps = 1e-16
idemTol = 1e-14

hDim = 98304 M = 80

Adjusted M = 96

Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 1 **local** row min = 6144 row max = 12288 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 3 **local** row min = 18432 row max = 24576 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 10 **local** row min = 61440 row max = 67584 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
total procs = 16 total rows = 98304 total cols = 96
global row min = 0 row max = 98304 row extent = 98304
rank = 0 **local** row min = 0 row max = 6144 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 5 **local** row min = 30720 row max = 36864 row extent = 6144

Sparsity:

Initial sparsity = 672042, fraction = 6.258879e-04, Avg per row = 6.836365
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 6 **local** row min = 36864 row max = 43008 row extent = 6144
Max per row = 7
l = 4, count = 2, fraction = 0.000020
l = 5, count = 621, fraction = 0.006317
l = 6, count = 14838, fraction = 0.150940
l = 7, count = 82843, fraction = 0.842723
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 13 **local** row min = 79872 row max = 86016 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 7 **local** row min = 43008 row max = 49152 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 12 **local** row min = 73728 row max = 79872 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 15 **local** row min = 92160 row max = 98304 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 8 **local** row min = 49152 row max = 55296 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 11 **local** row min = 67584 row max = 73728 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 14 **local** row min = 86016 row max = 92160 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 4 **local** row min = 24576 row max = 30720 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 9 **local** row min = 55296 row max = 61440 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 2 **local** row min = 12288 row max = 18432 row extent = 6144

Gershgorin:

New eMax, eMin = 1.745500e+00, -7.356212e-01
bufferSize = 9437184
Initial sparsity normalized = 672042, fraction = 6.258879e-04, avg = 6.83636, max = 7

SP2Loop:

iter = 0 trX = 4.935743e+04 trX2 = 2.720037e+04
iter = 1 trX = 2.720037e+04 trX2 = 9.994787e+03
iter = 2 trX = 4.440595e+04 trX2 = 2.485384e+04
iter = 3 trX = 6.395806e+04 trX2 = 4.735425e+04
iter = 4 trX = 4.735425e+04 trX2 = 3.149323e+04
iter = 5 trX = 6.321528e+04 trX2 = 5.026180e+04
iter = 6 trX = 5.026180e+04 trX2 = 3.881328e+04
iter = 7 trX = 3.881328e+04 trX2 = 2.922713e+04
iter = 8 trX = 4.839943e+04 trX2 = 4.062611e+04
iter = 9 trX = 5.617275e+04 trX2 = 4.981154e+04
iter = 10 trX = 4.981154e+04 trX2 = 4.464542e+04
iter = 11 trX = 4.464542e+04 trX2 = 4.032639e+04
iter = 12 trX = 4.896445e+04 trX2 = 4.554145e+04
iter = 13 trX = 5.238745e+04 trX2 = 4.956883e+04
iter = 14 trX = 4.956883e+04 trX2 = 4.731790e+04

mpirun has exited due to process rank 3 with PID 12638 on node e0700 exiting improperly. There are three reasons this could occur:

1. this process did not call "init" before exiting, but others in the job did. This can cause a job to hang indefinitely while it waits for all processes to call "init". By rule, if one process calls "init", then ALL processes must call "init" prior to termination.
2. this process called "init", but exited without calling "finalize". By rule, all processes that call "init" MUST call "finalize" prior to exiting or it will be considered an "abnormal termination"
3. this process called "MPI_Abort" or "orte_abort" and the mca parameter orte_create_session_dirs is set to false. In this case, the run-time cannot detect that the abort call was an abnormal termination. Hence, the only error message you will receive is this one.

This may have caused other processes in the application to be terminated by signals sent by mpirun (as reported here).

You can avoid this message by specifying -quiet on the mpirun command line.

CoSP2: SP2 Loop

Parameters:

msparse = 80 hDim = 98304 debug = 1
hmatName =
eps = 1e-05 hEps = 1e-16
idemTol = 1e-14

hDim = 98304 M = 80

Adjusted M = 96

Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 1 local row min = 6144 row max = 12288 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 3 local row min = 18432 row max = 24576 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 15 local row min = 92160 row max = 98304 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 4 local row min = 24576 row max = 30720 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 9 local row min = 55296 row max = 61440 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 13 local row min = 79872 row max = 86016 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 5 local row min = 30720 row max = 36864 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 11 local row min = 67584 row max = 73728 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 12 local row min = 73728 row max = 79872 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 8 local row min = 49152 row max = 55296 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 10 local row min = 61440 row max = 67584 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 6 local row min = 36864 row max = 43008 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 7 local row min = 43008 row max = 49152 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
rank = 14 local row min = 86016 row max = 92160 row extent = 6144
Generated H Matrix nnz = 672042 avg nnz/row = 6
total procs = 16 total rows = 98304 total cols = 96
global row min = 0 row max = 98304 row extent = 98304
rank = 0 local row min = 0 row max = 6144 row extent = 6144

Sparsity:

Initial sparsity = 672042, fraction = 6.258879e-04, Avg per row = 6.836365
Max per row = 7
l = 4, count = 2, fraction = 0.000020
l = 5, count = 621, fraction = 0.006317
l = 6, count = 14838, fraction = 0.150940
l = 7, count = 82843, fraction = 0.842723
Generated H Matrix nnz = 672042 avg nnz/row = 6

rank = 2 local row min = 12288 row max = 18432 row extent = 6144

Gershgorin:

New eMax, eMin = 1.745500e+00, -7.356212e-01

bufferSize = 9437184

Initial sparsity normalized = 672042, fraction = 6.258879e-04, avg = 6.83636, max = 7

SP2Loop:

iter = 10 trX = 4.981154e+04 trX2 = 4.464542e+04

iter = 11 trX = 4.464542e+04 trX2 = 4.032639e+04

iter = 12 trX = 4.896445e+04 trX2 = 4.554145e+04

iter = 13 trX = 5.238745e+04 trX2 = 4.956883e+04

iter = 14 trX = 4.956883e+04 trX2 = 4.731790e+04

iter = 15 trX = 4.731790e+04 trX2 = 4.544718e+04

iter = 16 trX = 4.918861e+04 trX2 = 4.771064e+04

iter = 17 trX = 4.771064e+04 trX2 = 4.649398e+04

iter = 18 trX = 4.892731e+04 trX2 = 4.795556e+04

iter = 19 trX = 4.989906e+04 trX2 = 4.910173e+04

iter = 20 trX = 4.910173e+04 trX2 = 4.855031e+04

iter = 21 trX = 4.965316e+04 trX2 = 5.060054e+04

iter = 22 trX = 4.870578e+04 trX2 = -9.750371e+05

iter = 23 trX = 1.072449e+06 trX2 = -5.136388e+12

iter = 24 trX = -5.136388e+12 trX2 = 7.295617e+24

Results:

X2 Sparsity CCN = 2906510, fraction = 2.706898e-03 avg = 29.5665, max = 89

D Sparsity AAN = 2906464, fraction = 2.706856e-03 avg = 29.5661, max = 89

Number of iterations = 25

Counters for Rank 0

Counter	Calls	Avg/Call(MB)	Total(MB)
reduce	19	0.0000	0.0003
send	29	2.6508	76.8721
recv	29	2.6315	76.3141

Counter Statistics Across 16 Ranks:

Counter	Rank: Min(MB)	Rank: Max(MB)	Avg(MB)	Stdev(MB)
reduce	0: 0.0003	0: 0.0003	0.0003	0.0000
send	15: 74.9711	7: 113.5838	106.5789	11.6620
recv	15: 75.9751	6: 112.0425	106.5789	11.5183

Timings for Rank 0

Timer	# Calls	Avg/Call (s)	Total (s)	% Loop
total	1	4.4862	4.4862	100.00
loop	1	4.4862	4.4862	100.00
pre	1	0.5449	0.5449	12.15
sp2Loop	1	3.7464	3.7464	83.51
norm	1	0.0439	0.0439	0.98
x2	15	0.0423	0.6340	14.13
xadd	8	0.1030	0.8236	18.36
xset	7	0.0369	0.2582	5.76
exchange	30	0.0033	0.0982	2.19
reduceComm	19	0.0210	0.3999	8.91

Timing Statistics Across 16 Ranks:

Timer	Rank: Min(s)	Rank: Max(s)	Avg(s)	Stdev(s)
total	1: 4.4668	10: 4.5643	4.5171	0.0276
loop	1: 4.4668	10: 4.5643	4.5171	0.0276
pre	1: 0.4197	2: 0.5793	0.5138	0.0399
sp2Loop	9: 3.7438	8: 3.7513	3.7490	0.0019
norm	1: 0.0081	6: 0.0463	0.0350	0.0112
x2	3: 0.1789	13: 0.7067	0.5970	0.1598
xadd	1: 0.0360	10: 0.8244	0.6514	0.2632
xset	3: 0.0240	7: 0.2859	0.2339	0.0798
exchange	0: 0.0982	3: 1.1791	0.4241	0.3105
reduceComm	7: 0.1875	1: 1.2948	0.4136	0.3341

LULESH

Livermore Unstructured Lagrangian Explicit Shock Hydrodynamics (LULESH)

File changes

In order to perform the cast from a C++ object to a char buffer, BOOST serialization was used. Three files were modified to port FTI: lulesh.cc, lulesh.h and lulesh-comm.cc. The modifications to the first two files are shown here. The modifications to the third file were barely the replacements of MPI_COMM_WORLD by FTI_COMM_WORLD and are not listed here.

```
diff --git a/LULESH/lulesh.cc b/FTI_LULESH/lulesh.cc
index a141611..d5572f8 100644
--- a/LULESH/lulesh.cc
+++ b/FTI_LULESH/lulesh.cc
@@ -162,6 +162,22 @@ Additional BSD Notice

#include "lulesh.h"

+//*****
+// Boost Serialization
+//*****
+#include <boost/archive/text_oarchive.hpp>
+#include <boost/archive/text_iarchive.hpp>
+
+#include <sstream>
+// --- File version ---
+#include <fstream>
+std::stringstream locDom_ser;
+
+//*****
+// FTI Checkpoint - Restart
+//*****
+#include <fti.h>
+#define ITER_CKPT 500

/*****
/* Data structure implementation */
@@ -213,7 +229,7 @@ void TimeIncrement(Domain& domain)
#if USE_MPI
    MPI_Allreduce(&gnewdt, &newdt, 1,
        ((sizeof(Real_t) == 4) ? MPI_FLOAT : MPI_DOUBLE),
-        MPI_MIN, MPI_COMM_WORLD);
+        MPI_MIN, FTI_COMM_WORLD);
#else
    newdt = gnewdt;
#endif
@@ -1061,7 +1077,7 @@ void CalcHourglassControlForElems(Domain& domain,
/* Do a check for negative volumes */
    if ( domain.v(i) <= Real_t(0.0) ) {
#if USE_MPI
-        MPI_Abort(MPI_COMM_WORLD, VolumeError);
+        MPI_Abort(FTI_COMM_WORLD, VolumeError);
#else
        exit(VolumeError);
#endif
@@ -1111,7 +1127,7 @@ void CalcVolumeForceForElems(Domain& domain)
    for ( Index_t k=0 ; k<numElem ; ++k ) {
        if (determ[k] <= Real_t(0.0)) {
#if USE_MPI
-        MPI_Abort(MPI_COMM_WORLD, VolumeError);
+        MPI_Abort(FTI_COMM_WORLD, VolumeError);
#else
        exit(VolumeError);
#endif
@@ -1626,7 +1642,7 @@ void CalcLagrangeElements(Domain& domain, Real_t* vnew)
    if (vnew[k] <= Real_t(0.0))
    {
#if USE_MPI
```

```

- MPI_Abort(MPI_COMM_WORLD, VolumeError) ;
+ MPI_Abort(FTI_COMM_WORLD, VolumeError) ;
#else
    exit(VolumeError);
#endif
@@ -2030,7 +2046,7 @@ void CalcQForElems(Domain& domain, Real_t vnew[])

    if(idx >= 0) {
    #if USE_MPI
- MPI_Abort(MPI_COMM_WORLD, QStopError) ;
+ MPI_Abort(FTI_COMM_WORLD, QStopError) ;
    #else
        exit(QStopError);
    #endif
    @@ -2399,7 +2415,7 @@ void ApplyMaterialPropertiesForElems(Domain& domain, Real_t vnew[])
    }
    if (vc <= 0.) {
    #if USE_MPI
- MPI_Abort(MPI_COMM_WORLD, VolumeError) ;
+ MPI_Abort(FTI_COMM_WORLD, VolumeError) ;
    #else
        exit(VolumeError);
    #endif
    @@ -2683,6 +2699,19 @@ void LagrangeLeapFrog(Domain& domain)
    #endif
    }

+//Serialization
+void save (Domain *dom_saved){
+ boost::archive::text_oarchive oa(locDom_ser);
+ oa << dom_saved;
+}
+
+//Deserialization
+Domain* load (){
+ Domain *dom_loaded;
+ boost::archive::text_iarchive ia(locDom_ser);
+ ia >> dom_loaded;
+ return dom_loaded;
+}

/*****/

@@ -2697,8 +2726,10 @@ int main(int argc, char *argv[])
    Domain_member fieldData ;

    MPI_Init(&argc, &argv) ;
- MPI_Comm_size(MPI_COMM_WORLD, &numRanks) ;
- MPI_Comm_rank(MPI_COMM_WORLD, &myRank) ;
+ char config_fti[] = "config.fti";
+ FTI_Init(config_fti, MPI_COMM_WORLD);
+ MPI_Comm_size(FTI_COMM_WORLD, &numRanks) ;
+ MPI_Comm_rank(FTI_COMM_WORLD, &myRank) ;
    #else
        numRanks = 1;
        myRank = 0;
    @@ -2755,7 +2786,7 @@ int main(int argc, char *argv[])
        CommSBN(*locDom, 1, &fieldData) ;

        // End initialization
- MPI_Barrier(MPI_COMM_WORLD);
+ MPI_Barrier(FTI_COMM_WORLD);
    #endif

    // BEGIN timestep to solution */
    @@ -2766,10 +2797,68 @@ int main(int argc, char *argv[])
        gettimeofday(&start, NULL) ;
    #endif
    //debug to see region sizes
-// for(Int_t i = 0; i < locDom->numReg(); i++)
-// std::cout << "region" << i + 1 << "size" << locDom->regElemSize(i) <<std::endl;
- while((locDom->time() < locDom->stoptime()) && (locDom->cycle() < opts.its)) {
+ // for(Int_t i = 0; i < locDom->numReg(); i++)
+ // std::cout << "region" << i + 1 << "size" << locDom->regElemSize(i) <<std::endl;

```

```

+
+
+//-----//
+
+ //First serialization to get a buffer size
+ save(locDom);

+ //Cast std::stringstream -> char*
+ int buffer_size = 0;
+ char* buffer_locDom_ser;
+ std::string tmp = locDom_ser.str();
+ buffer_size = tmp.size();
+ buffer_size += 1000000; //Add this to handle the dynamic change size of the buffer
+ buffer_locDom_ser = new char [buffer_size];
+ strcpy(buffer_locDom_ser, tmp.c_str());
+
+ //Checkpoint informations
+ int id = 1;
+ int level = 1;
+ int res;
+
+ FTI_Protect(0, &id, 1, FTI_INTG);
+ FTI_Protect(1, &level, 1, FTI_INTG);
+ FTI_Protect(2, buffer_locDom_ser, buffer_size, FTI_CHAR);
+
+
+ //Restart
+ if(FTI_Status() != 0){
+   if(!myRank)
+     std::cout << "---- Restart ----\n";
+   res = FTI_Recover();
+
+   //Update checkpoint information
+   if (res != 0) {
+     exit(1);
+   }
+   else { // Update ckpt. id & level
+     level = (level+1)%5;
+     id++;
+   }
+
+   //Cast char* to stringstream
+   locDom_ser.str(""); //reset the stringstream
+   locDom_ser.str(buffer_locDom_ser);
+
+   //Deserialization
+   Domain *tmp;
+   tmp = load();
+
+   //Set the used by simulation object
+   delete locDom;
+   locDom = NULL;
+   locDom = tmp;
+ }
+
+//-----//
+ if (!myRank)
+   std::cout << "-- Start of the main loop --\n";
+ while((locDom->time() < locDom->stoptime()) && (locDom->cycle() < opts.its)) {
+   TimeIncrement(*locDom);
+   LagrangeLeapFrog(*locDom);

@@ -2777,6 +2866,26 @@ int main(int argc, char *argv[])
    printf("cycle = %d, time = %e, dt=%e\n",
          locDom->cycle(), double(locDom->time()), double(locDom->deltatime()) );
  }

+ //Checkpoint at ITER_CKPT
+ if((locDom->cycle()%ITER_CKPT) == 0 && locDom->cycle() != opts.its){
+
+   //Serialization of locDom in std::stringstream
+   locDom_ser.str("");
+   save(locDom);
+
+

```



```

+ //Cast std::stringstream -> char*
+ std::string tmp = locDom_ser.str();
+ buffer_locDom_ser[0] = '\0'; //reset the buffer
+ strcpy(buffer_locDom_ser, tmp.c_str());
+
+ res = FTI_Checkpoint(id, level);
+ // sleep(3); //for the tests
+ if(res != 0){
+     id++;
+     level= (level%4)+1;
+ }
+ }
+ }
}

// Use reduced max elapsed time
@@ -2791,7 +2900,7 @@ int main(int argc, char *argv[])
    double elapsed_timeG;
    #if USE_MPI
        MPI_Reduce(&elapsed_time, &elapsed_timeG, 1, MPI_DOUBLE,
-         MPI_MAX, 0, MPI_COMM_WORLD);
+         MPI_MAX, 0, FTI_COMM_WORLD);
    #else
        elapsed_timeG = elapsed_time;
    #endif
@@ -2806,6 +2915,7 @@ int main(int argc, char *argv[])
}

#if USE_MPI
+ FTI_Finalize();
    MPI_Finalize();
#endif

```

```

diff --git a/LULESH/lulesh.h b/FTI_LULESH/lulesh.h
index b6afd5c..1ca6a59 100644
--- a/LULESH/lulesh.h
+++ b/FTI_LULESH/lulesh.h
@@ -24,6 +24,16 @@
#include <math.h>
#include <vector>

+//*****
+// Boost Serialization
+//*****
+#include <boost/serialization/vector.hpp>
+#include <iostream>
+#include <fstream>
+#if _OPENMP
+#include <omp.h>
+#endif
+
+//*****
+// Allow flexibility for arithmetic representations
+//*****
@@ -133,6 +143,27 @@ class Domain {
    Index_t rowLoc, Index_t planeLoc,
    Index_t nx, Int_t tp, Int_t nr, Int_t balance, Int_t cost);

+ Domain () :
+ m_e_cut(Real_t(1.0e-7)),
+ m_p_cut(Real_t(1.0e-7)),
+ m_q_cut(Real_t(1.0e-7)),
+ m_v_cut(Real_t(1.0e-10)),
+ m_u_cut(Real_t(1.0e-7)),
+ m_hgcoef(Real_t(3.0)),
+ m_ss4o3(Real_t(4.0)/Real_t(3.0)),
+ m_qstop(Real_t(1.0e+12)),
+ m_monoq_max_slope(Real_t(1.0)),
+ m_monoq_limiter_mult(Real_t(2.0)),
+ m_qlc_monoq(Real_t(0.5)),
+ m_qqc_monoq(Real_t(2.0)/Real_t(3.0)),
+ m_qqc(Real_t(2.0)),
+ m_eosvmax(Real_t(1.0e+9)),
+ m_eosvmin(Real_t(1.0e-9))

```

```

+ m_eccentricity(Real_t(1.0e-9));
+ m_pmin(Real_t(0.)),
+ m_emin(Real_t(-1.0e+15)),
+ m_dvovmax(Real_t(0.1)),
+ m_refdens(Real_t(1.0)) {};
+
+ //
+ // ALLOCATION
+ //
@@ -423,6 +454,243 @@ class Domain {
    void SetupElementConnectivities(Int_t edgeElems);
    void SetupBoundaryConditions(Int_t edgeElems);

+ friend class boost::serialization::access;
+ template <typename Archive>
+ void serialize(Archive &ar, const unsigned int version){
+
+ //Check de/serialization
+ // if(Archive::is_loading::value){
+ //   std::cout << "-----\n";
+ //   std::cout << "Start of deserialization.\n";
+ //   std::cout << "-----\n";
+ // }
+ // else {
+ //   std::cout << "-----\n";
+ //   std::cout << "Start of serialization.\n";
+ //   std::cout << "-----\n";
+ // }
+
+ ar & m_x ; /* coordinates */
+ ar & m_y;
+ ar & m_z;
+
+ ar & m_xd ; /* velocities */
+ ar & m_yd ;
+ ar & m_zd ;
+
+ ar & m_xdd ; /* accelerations */
+ ar & m_ydd ;
+ ar & m_zdd ;
+
+ ar & m_fx ; /* forces */
+ ar & m_fy ;
+ ar & m_fz ;
+
+ ar & m_nodalMass ; /* mass */
+
+ ar & m_symmX ; /* symmetry plane nodesets */
+ ar & m_symmY ;
+ ar & m_symmZ ;
+
+ // Element-centered
+
+ ar & m_numRanks ;
+ ar & m_colLoc ;
+ ar & m_rowLoc ;
+ ar & m_planeLoc ;
+ ar & m_tp ;
+
+ ar & m_sizeX ;
+ ar & m_sizeY ;
+ ar & m_sizeZ ;
+ ar & m_numElem ;
+ ar & m_numNode ;
+
+ ar & m_maxPlaneSize ;
+ ar & m_maxEdgeSize ;
+
+ // Region information
+ ar & m_numReg ;
+ ar & m_cost; //imbalance cost
+
+ if(Archive::is_loading::value){
+   m_regElemSize = new Index_t[m_numReg];

```

```

+ }
+ ar & boost::serialization::make_array <Index_t> (m_regElemSize, m_numReg); // Size of region sets
+
+ if(Archive::is_loading::value){
+     m_regNumList = new Index_t[m_numElem];
+ }
+ ar & boost::serialization::make_array <Index_t> (m_regNumList, m_numElem); // Region number per domain element
+
+ if(Archive::is_loading::value){
+     m_regElemList = new Index_t*[m_numReg];
+     for (int i = 0; i < m_numReg; i++){
+         m_regElemList[i] = new Index_t[m_regElemSize[i]];
+     }
+ }
+
+ for (int i = 0; i < m_numReg; i++){
+     ar & boost::serialization::make_array <Index_t> (m_regElemList[i], m_regElemSize[i]);
+ }
+
+ ar & m_nodelist ; /* elemToNode connectivity */
+
+ ar & m_lxim ; /* element connectivity across each face */
+ ar & m_lxip ;
+ ar & m_letam ;
+ ar & m_letap ;
+ ar & m_lzetam ;
+ ar & m_lzetap ;
+
+ ar & m_elemBC ; /* symmetry/free-surface flags for each elem face */
+
+ ar & m_dxx ; /* principal strains -- temporary */
+ ar & m_dyy ;
+ ar & m_dzz ;
+
+ ar & m_delv_xi ; /* velocity gradient -- temporary */
+ ar & m_delv_eta ;
+ ar & m_delv_zeta ;
+
+ ar & m_delx_xi ; /* coordinate gradient -- temporary */
+ ar & m_delx_eta ;
+ ar & m_delx_zeta ;
+
+ ar & m_e ; /* energy */
+
+ ar & m_p ; /* pressure */
+ ar & m_q ; /* q */
+ ar & m_ql ; /* linear term for q */
+ ar & m_qq ; /* quadratic term for q */
+
+ ar & m_v ; /* relative volume */
+ ar & m_volo ; /* reference volume */
+ ar & m_vnew ; /* new relative volume -- temporary */
+ ar & m_delv ; /* m_vnew - m_v */
+ ar & m_vdov ; /* volume derivative over volume */
+
+ ar & m_arealg ; /* characteristic length of an element */
+
+ ar & m_ss ; /* "sound speed" */
+
+ ar & m_elemMass ; /* mass */
+
+ // Cutoffs (treat as constants)
+ ar & const_cast<Real_t &>(m_e_cut);
+ ar & const_cast<Real_t &>(m_p_cut);
+ ar & const_cast<Real_t &>(m_q_cut);
+ ar & const_cast<Real_t &>(m_v_cut);
+ ar & const_cast<Real_t &>(m_u_cut);
+
+ // Other constants (usually settable, but hardcoded in this proxy app)
+ ar & const_cast<Real_t &>(m_hgcoef);
+ ar & const_cast<Real_t &>(m_ss4o3);
+ ar & const_cast<Real_t &>(m_qstop);
+ ar & const_cast<Real_t &>(m_monoq_max_slope);
+ ar & const_cast<Real_t &>(m_monoq_limiter_mult);
+ ar & const_cast<Real_t &>(m_elo_monoc);

```

```

+   ar & const_cast<Real_t &>(m_qc_monoq);
+   ar & const_cast<Real_t &>(m_qqc_monoq);
+   ar & const_cast<Real_t &>(m_qqc);
+   ar & const_cast<Real_t &>(m_eosvmax);
+   ar & const_cast<Real_t &>(m_eosvmin);
+   ar & const_cast<Real_t &>(m_pmin);
+   ar & const_cast<Real_t &>(m_emin);
+   ar & const_cast<Real_t &>(m_dvovmax);
+   ar & const_cast<Real_t &>(m_refdens);
+
+   // Variables to keep track of timestep, simulation time, and cycle
+   ar & m_dtcourant ; // courant constraint
+   ar & m_dthydro ; // volume change constraint
+   ar & m_cycle ; // iteration count for simulation
+   ar & m_dtfixed ; // fixed time increment
+   ar & m_time ; // current time
+   ar & m_deltatime ; // variable time increment
+   ar & m_deltatimemultlb ;
+   ar & m_deltatimemultub ;
+   ar & m_dtmax ; // maximum allowable time increment
+   ar & m_stoptime ; // end time for simulation
+
+   // OMP hack
+   #if _OPENMP
+       Index_t numthreads = omp_get_max_threads();
+   #else
+       Index_t numthreads = 1;
+   #endif
+
+   if (numthreads > 1) {
+       if(Archive::is_loading::value){
+           m_nodeElemStart = new Index_t[m_numNode+1];
+       }
+       ar & boost::serialization::make_array <Index_t> (m_nodeElemStart, m_numNode+1);
+
+       if(Archive::is_loading::value){
+           m_nodeElemCornerList = new Index_t[m_nodeElemStart[m_numNode]];
+       }
+       ar & boost::serialization::make_array <Index_t> (m_nodeElemCornerList, m_nodeElemStart[m_numNode]);
+   } else {
+       m_nodeElemStart = NULL;
+       m_nodeElemCornerList = NULL;
+   }
+
+   // Used in setup
+   ar & m_rowMin;
+   ar & m_rowMax;
+   ar & m_colMin;
+   ar & m_colMax;
+   ar & m_planeMin;
+   ar & m_planeMax;
+
+   #if USE_MPI
+       // account for face communication
+       Index_t comBufSize =
+           (m_rowMin + m_rowMax + m_colMin + m_colMax + m_planeMin + m_planeMax) *
+           m_maxPlaneSize * MAX_FIELDS_PER_MPI_COMM ;
+
+       // account for edge communication
+       comBufSize +=
+           ((m_rowMin & m_colMin) + (m_rowMin & m_planeMin) + (m_colMin & m_planeMin) +
+            (m_rowMax & m_colMax) + (m_rowMax & m_planeMax) + (m_colMax & m_planeMax) +
+            (m_rowMax & m_colMin) + (m_rowMin & m_planeMax) + (m_colMin & m_planeMax) +
+            (m_rowMin & m_colMax) + (m_rowMax & m_planeMin) + (m_colMax & m_planeMin)) *
+           m_maxEdgeSize * MAX_FIELDS_PER_MPI_COMM ;
+
+       // account for corner communication
+       // factor of 16 is so each buffer has its own cache line
+       comBufSize += ((m_rowMin & m_colMin & m_planeMin) +
+           (m_rowMin & m_colMin & m_planeMax) +
+           (m_rowMin & m_colMax & m_planeMin) +
+           (m_rowMin & m_colMax & m_planeMax) +
+           (m_rowMax & m_colMin & m_planeMin) +
+           (m_rowMax & m_colMin & m_planeMax) +
+           (m_rowMax & m_colMax & m_planeMin) +

```

```

+      (m_rowMax & m_colMax & m_planeMax)) * CACHE_COHERENCE_PAD_REAL ;
+
+
+  // Communication Work space
+  if(Archive::is_loading::value){
+    commDataSend = new Real_t[comBufSize];
+    commDataRecv = new Real_t[comBufSize];
+  }
+  ar & boost::serialization::make_array <Real_t> (commDataRecv,comBufSize);
+  ar & boost::serialization::make_array <Real_t> (commDataSend,comBufSize);
+
+  #endif
+
+  //Check de/serialization
+  // if(Archive::is_loading::value){
+  //   std::cout << "-----\n";
+  //   std::cout << "Deserialization finished.\n";
+  //   std::cout << "-----\n";
+  // }
+  // else {
+  //   std::cout << "-----\n";
+  //   std::cout << "Serialization finished.\n";
+  //   std::cout << "-----\n";
+  // }
+ }
+
+  //
+  // IMPLEMENTATION
+  //

```